



# **Natural Language Processing**

## **for Music**

**Sergio Oramas (Pandora)**



# Outline

## Section I: Natural Language Processing

1. Introduction to NLP
2. Semantics
3. Text Representation
4. Deep Learning Architectures
5. Contextualized Embeddings
6. Multimodality
7. Conversational Systems

# Outline

## Section II: Use cases in Music

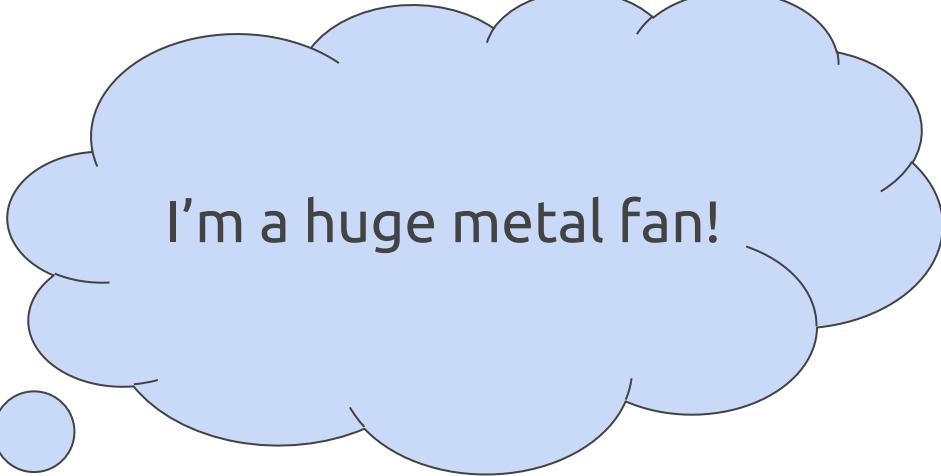
1. Automatic Knowledge Base Construction
2. Music Recommendation with Knowledge Graphs
3. Multimodal Music Recommendation
4. Multimodal Music Genre Classification
5. NLP for Music Knowledge Discovery

# Introduction to NLP

# What is Natural Language Processing?

- NLP is a field of Computer Science and Artificial Intelligence concerned with the interaction between computers and human (natural) language.
- Alan Turing's paper *Computing Machinery and Intelligence* is believed to be the first NLP paper. It stated that a computer could be considered intelligent if it could carry on a conversation with a human being without the human realizing he/she were talking to a machine.
- NLP is a core component in daily life technologies: web search, speech recognition and synthesis, automatic summaries in the web, product (including music) recommendation, machine translation...

# Why is it hard?



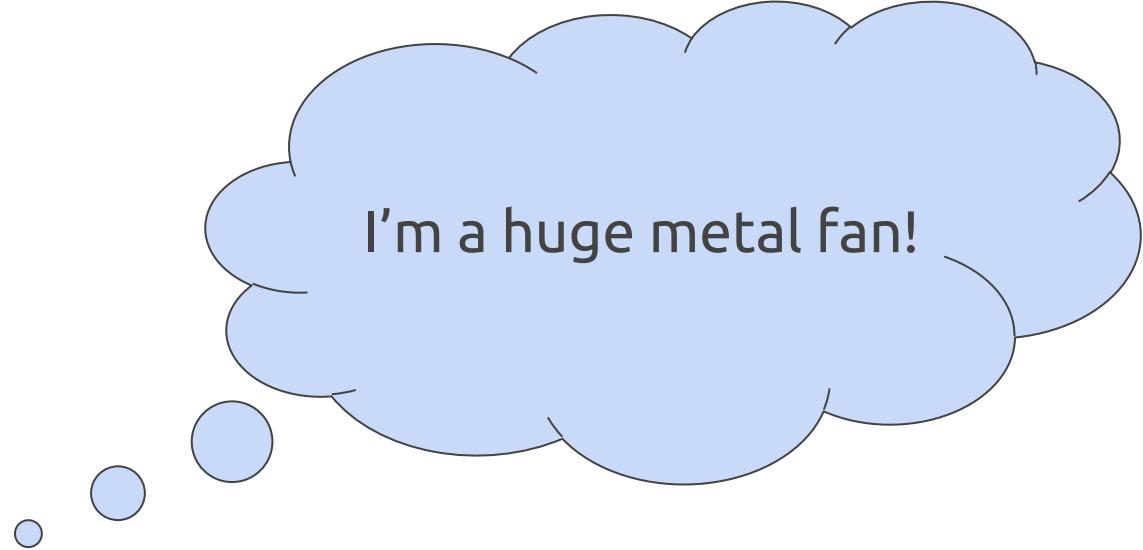
I'm a huge metal fan!

# Why is it hard?

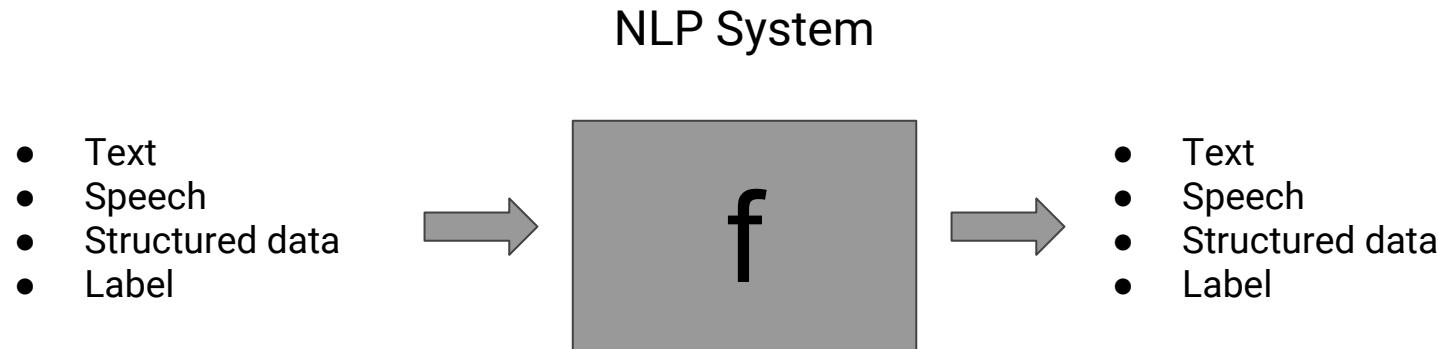


I'm a huge metal fan!

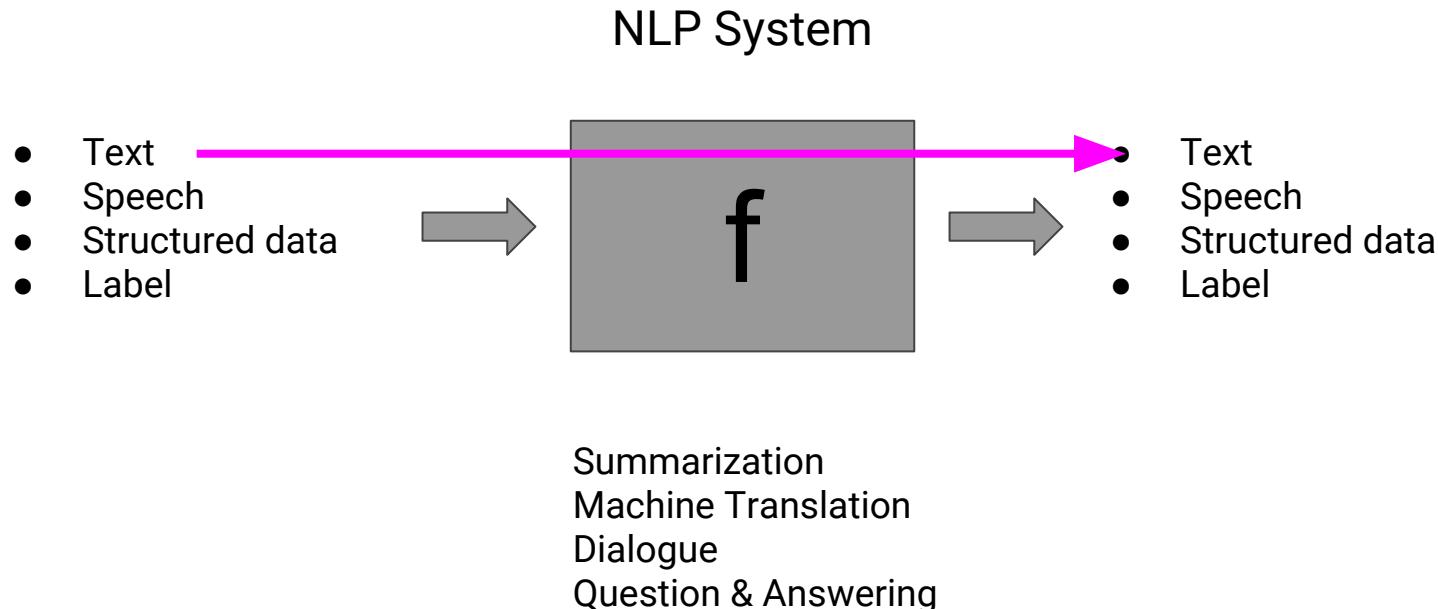
# Why is it hard?



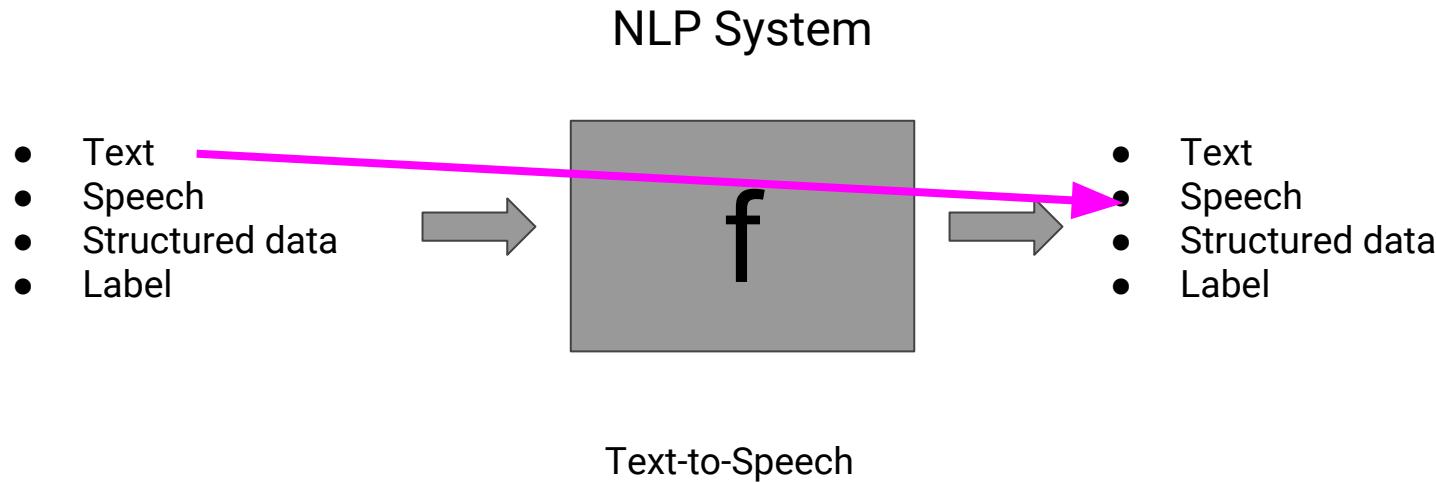
# Natural Language Processing



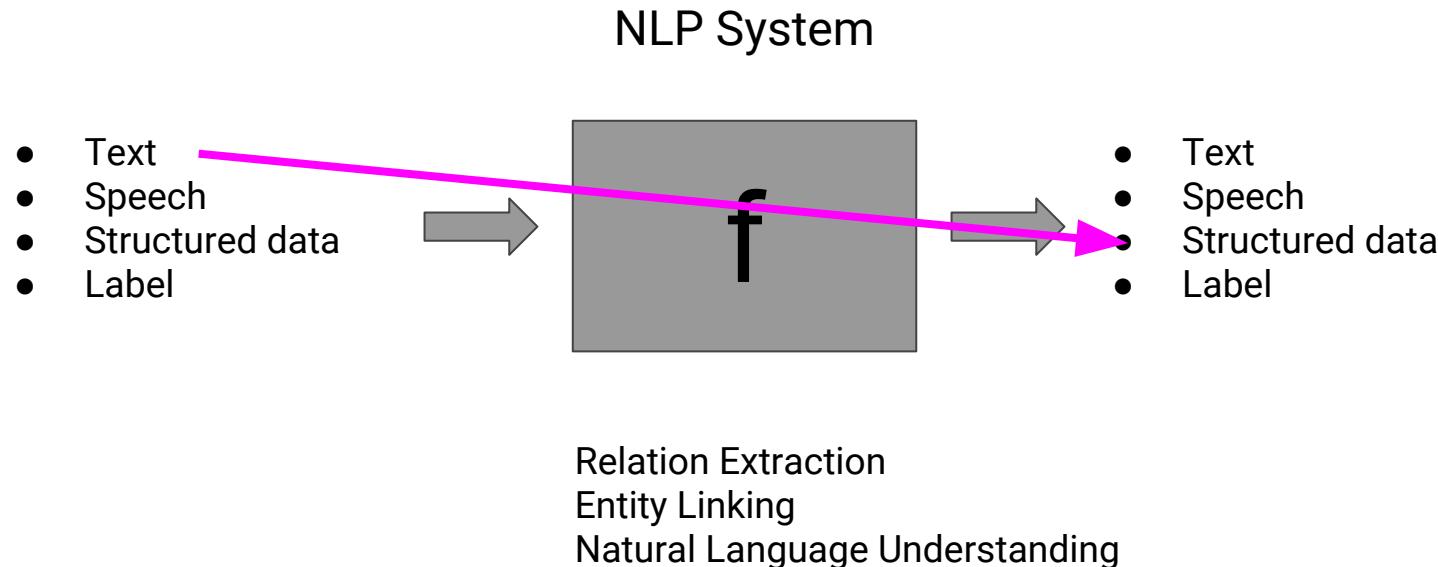
# Natural Language Processing Tasks



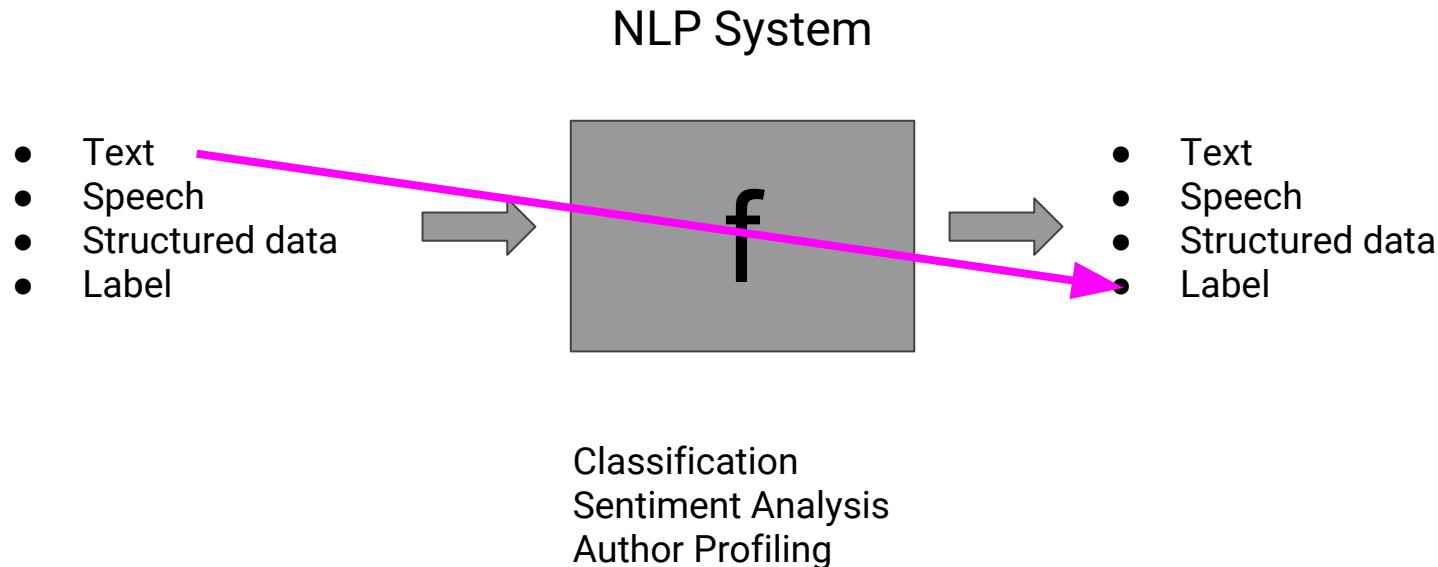
# Natural Language Processing Tasks



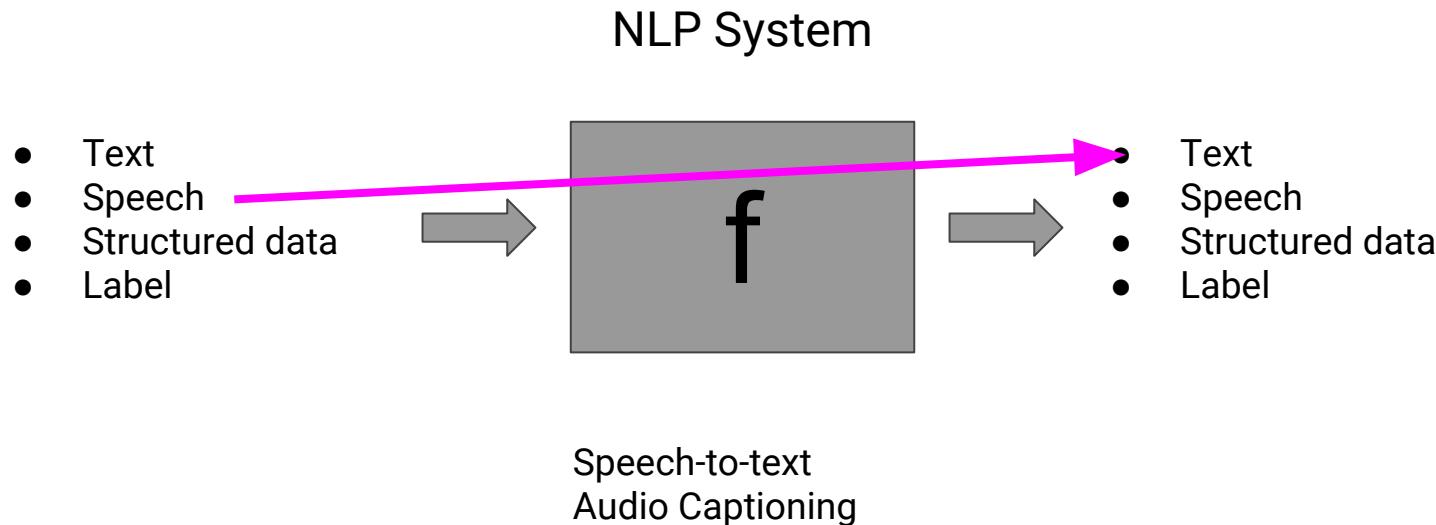
# Natural Language Processing Tasks



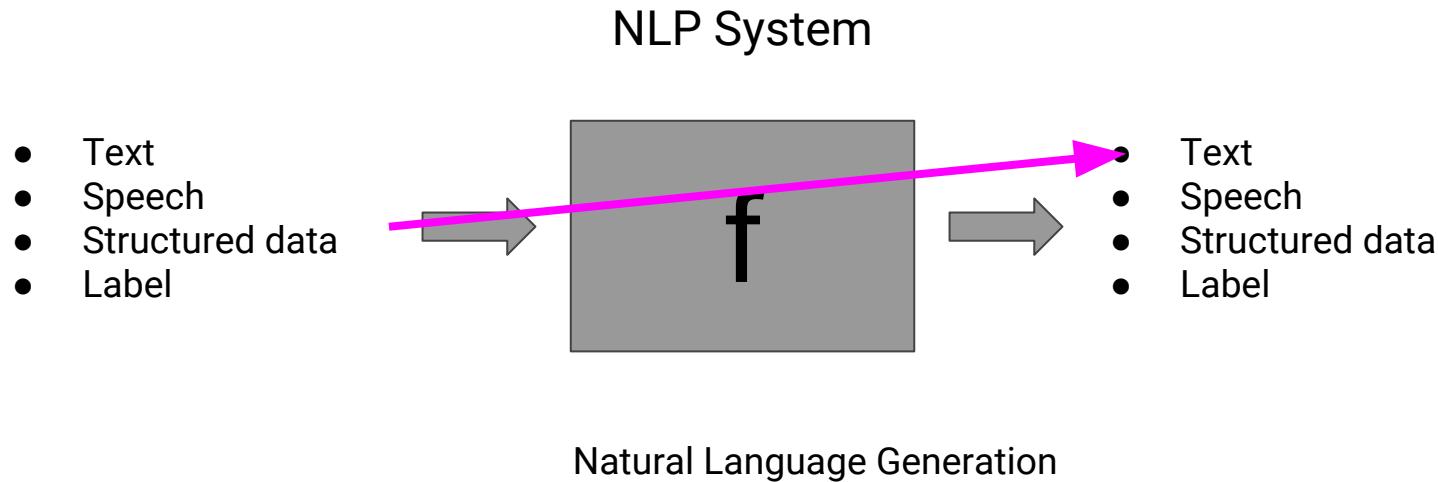
# Natural Language Processing Tasks



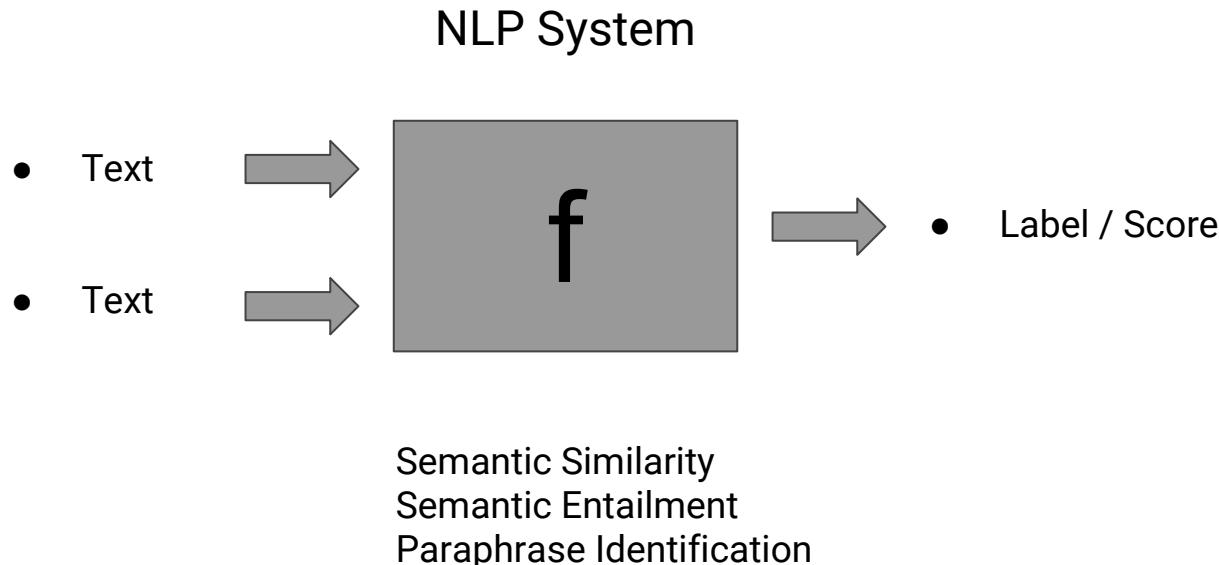
# Natural Language Processing Tasks



# Natural Language Processing Tasks



# Natural Language Processing Tasks



# Core Linguistic NLP Tasks

- \* Part-of-speech Tagging
- \* Syntactic Parsing
- \* Semantic Parsing
- \* Named Entity Recognition
- \* Coreference Resolution
- \* Word Sense Disambiguation

# *Core* elements in NLP - Part-of-Speech Tagging

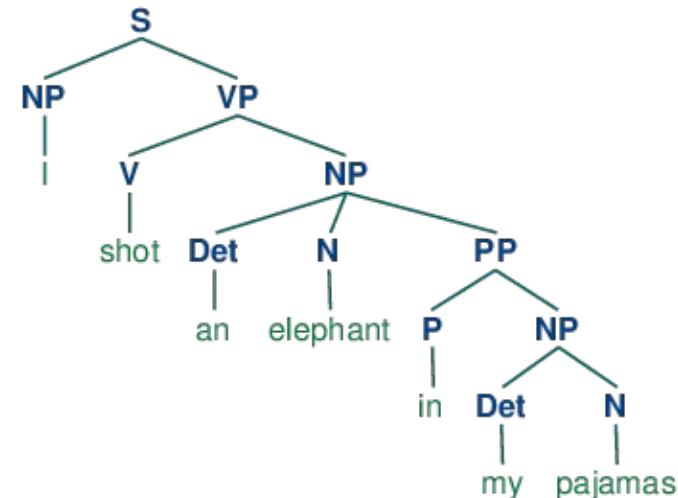
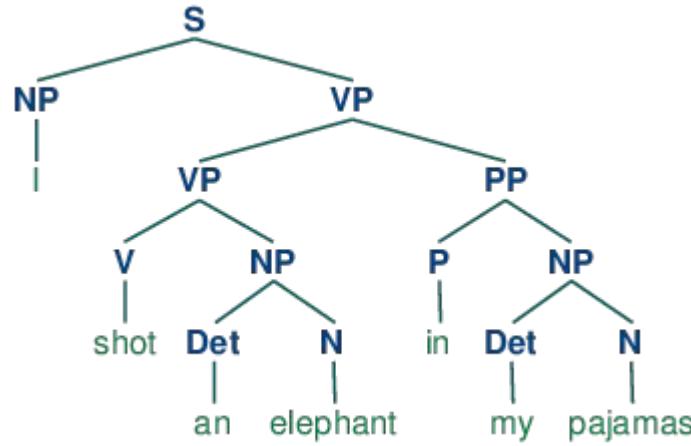
I like jazz music, it's like being alive for a second.

# Core elements in NLP - Part-of-Speech Tagging



# Core elements in NLP - Syntactic Parsing

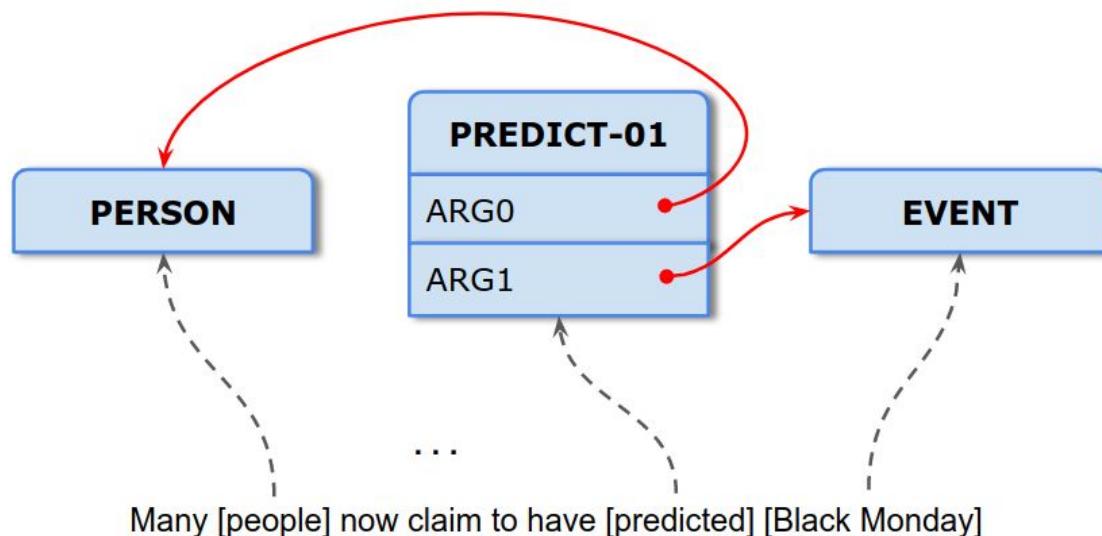
Identify relations holding between words or phrases in the sentence, and what is their *function*.



<http://www.nltk.org/book/ch08.html>

# Core elements in NLP - Semantic Parsing

A level of parsing above morphology and syntax. Capture underlying semantics expressed in language. Most focus on verbs and their *arguments*.



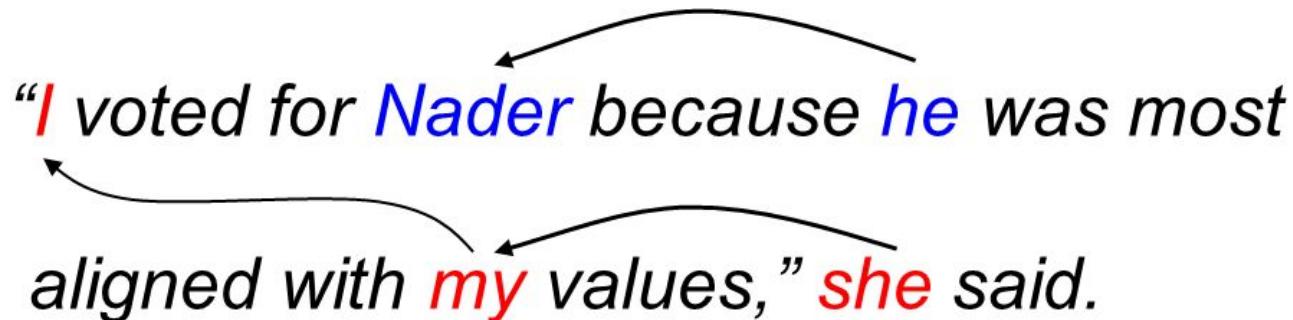
# Core elements in NLP - Named Entity Recognition

• Manfred Mann's Earth Band is a British progressive rock group formed in 1971 by Manfred Mann, a keyboard player born in South Africa best known as a founding member and namesake of 60s group Manfred Mann.

Band  
Music Genre  
Artist  
Country

# Core elements in NLP - Coreference Resolution

*“I voted for Nader because he was most aligned with my values,” she said.*



The diagram illustrates coreference resolution with three arrows pointing from pronouns to their antecedents: one from "he" to "Nader", one from "she" to "she" (the second occurrence), and one from "my" to "she".

<http://nlp.stanford.edu/projects/coref.shtml>

# *Core* elements in NLP - Word Sense Disambiguation

- “The performance of that **bass** player was outstanding”

# *Core* elements in NLP - Word Sense Disambiguation

- “The performance of that **bass** player was outstanding”



<https://tackyraocoon.com/2011/11/21/all-your-bass-are-belong-to-us/>

# Tools

Alchemy API

<http://www.alchemyapi.com/products/alchemylanguage/entity-extraction>

AYLIEN API <http://aylien.com/text-api>

Stanford NLP <http://nlp.stanford.edu:8080/sentiment/rntnDemo.html>

Gensim python library <https://radimrehurek.com/gensim/>

Senti WordNet <http://sentiwordnet.isti.cnr.it/>

# Software

## Standalone

- OpenNLP: <https://opennlp.apache.org/>
- Stanford CoreNLP: <http://stanfordnlp.github.io/CoreNLP/>
- Freeling: <http://nlp.lsi.upc.edu/freeling/node/1>
- Gate: <https://gate.ac.uk/>
- Mate Parser:  
<http://www.ims.uni-stuttgart.de/forschung/ressourcen/werkzeuge/matetools.en.html>

## Python Libraries

- Spacy: <https://spacy.io>
- Pattern: <http://www.clips.ua.ac.be/pattern>
- NLTK: <http://www.nltk.org/>
- Gensim: <https://radimrehurek.com/gensim/>
- Blob: <http://textblob.readthedocs.io/en/dev/>
- Rake:  
<https://www.airpair.com/nlp/keyword-extraction-tutorial>

# Software

## ML toolkits/libraries widely used in NLP

- CRF++: <https://taku910.github.io/crfpp/>
- Mallet: <http://mallet.cs.umass.edu/>
- Networkx: <https://networkx.github.io>
- Weka: <http://www.cs.waikato.ac.nz/ml/weka/>
- Deep Learning:
  - Keras <https://keras.io/>
  - Tflearn <http://tflearn.org/>
  - Tensorflow <https://www.tensorflow.org/>
  - Theano <http://deeplearning.net/software/theano/>
  - DyNet (formerly cnn) <https://github.com/clab/dynet>

# References - NLP

**Part-of-Speech Tagging:** Schmid, H. (1994, September). Probabilistic part-of-speech tagging using decision trees. In Proceedings of the international conference on new methods in language processing (Vol. 12, pp. 44-49).

**Parsing:** Chomsky, N. (2002). Syntactic structures. Walter de Gruyter.; Nivre, J. (2003). An efficient algorithm for projective dependency parsing. In Proceedings of the 8th International Workshop on Parsing Technologies (IWPT).

**Named Entity Recognition:** Tjong Kim Sang, E. F., & De Meulder, F. (2003, May). Introduction to the CoNLL-2003 shared task: Language-independent named entity recognition. In Proceedings of the seventh conference on Natural language learning at HLT-NAACL 2003-Volume 4 (pp. 142-147). Association for Computational Linguistics.

**Coreference Resolution:** Soon, W. M., Ng, H. T., & Lim, D. C. Y. (2001). A machine learning approach to coreference resolution of noun phrases. Computational linguistics, 27(4), 521-544.

**Summarization:** Saggion, H., & Lapalme, G. (2002). Generating indicative-informative summaries with sumUM. Computational linguistics, 28(4), 497-526.

**Simplification:** Chandrasekar, R., Doran, C., & Srivivas, B. (1996, August). Motivations and methods for text simplification. In Proceedings of the 16th conference on Computational linguistics-Volume 2 (pp. 1041-1044). Association for Computational Linguistics.

**Sentiment Analysis:** Pang, B., & Lee, L. (2008). Opinion mining and sentiment analysis. Foundations and trends in information retrieval, 2(1-2), 1-135.

**Author Profiling:** Estival, D., Gaustad, T., Pham, S. B., Radford, W., & Hutchinson, B. (2007). Author profiling for English emails. In Proceedings of the 10th Conference of the Pacific Association for Computational Linguistics (PACLING'07) (pp. 263-272).

**Topic Modeling:** Wallach, H. M. (2006, June). Topic modeling: beyond bag-of-words. In Proceedings of the 23rd international conference on Machine learning (pp. 977-984). ACM.

**Machine Translation:** Koehn, P., Hoang, H., Birch, A., Callison-Burch, C., Federico, M., Bertoldi, N., ... & Dyer, C. (2007, June). Moses: Open source toolkit for statistical machine translation. In Proceedings of the 45th annual meeting of the ACL on interactive poster and demonstration sessions (pp. 177-180). Association for Computational Linguistics.

**Lexical Semantics:** Cruse, D. A. (1986). Lexical semantics. Cambridge University Press.

**Word Sense Disambiguation.** Navigli, R. (2009). "Word sense disambiguation: A survey." ACM Computing Surveys (CSUR) 41.2: 10.

# References - KBs

**WordNet:** Miller, George A. "WordNet: a lexical database for English." *Communications of the ACM* 38.11 (1995): 39-41.

**Chebi:** Degtarenko, Kirill, et al. "ChEBI: a database and ontology for chemical entities of biological interest." *Nucleic acids research* 36.suppl 1 (2008): D344-D350.

**Snomed:** Spackman, Kent A., Keith E. Campbell, and Roger A. Côté. "SNOMED RT: a reference terminology for health care." *Proceedings of the AMIA annual fall symposium*. American Medical Informatics Association, 1997.

**BabelNet:** Navigli, Roberto, and Simone Paolo Ponzetto. "BabelNet: The automatic construction, evaluation and application of a wide-coverage multilingual semantic network." *Artificial Intelligence* 193 (2012): 217-250.

**DBpedia:** Auer, Sören, et al. "Dbpedia: A nucleus for a web of open data." *The semantic web*. Springer Berlin Heidelberg, 2007. 722-735.

**Yago:** Suchanek, Fabian M., Gjergji Kasneci, and Gerhard Weikum. "Yago: a core of semantic knowledge." *Proceedings of the 16th international conference on World Wide Web*. ACM, 2007.

**NELL:** Carlson, Andrew, et al. "Toward an Architecture for Never-Ending Language Learning." *AAAI*. Vol. 5. 2010.

**PATTY:** Nakashole, Ndapandula, Gerhard Weikum, and Fabian Suchanek. "PATTY: a taxonomy of relational patterns with semantic types." *Proceedings of the 2012 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*. Association for Computational Linguistics, 2012.

**WiSeNet:** Moro, Andrea, and Roberto Navigli. "WiSeNet: Building a Wikipedia-based semantic network with ontologized relations." *Proceedings of the 21st ACM international conference on Information and knowledge management*. ACM, 2012.

**DefIE:** Delli Bovi, Claudio, Luca Telesca, and Roberto Navigli. "Large-Scale Information Extraction from Textual Definitions through Deep Syntactic and Semantic Analysis." *Transactions of the Association for Computational Linguistics* 3 (2015): 529-543.

**KB-Unify:** Bovi, Claudio Delli, Luis Espinosa Anke, and Roberto Navigli. "Knowledge Base Unification via Sense Embeddings and Disambiguation." *Proceedings of EMNLP*. 2015.

**MusicBrainz:** Swartz, Aaron. "Musicbrainz: A semantic web service." *IEEE Intelligent Systems* 17.1 (2002): 76-77.

**Discogs:** [www.discogs.com](http://www.discogs.com)

**Grove Online:** [http://www.oxfordmusiconline.com/public/book/omo\\_gmo](http://www.oxfordmusiconline.com/public/book/omo_gmo)

# References

- Pang, B., & Lee, L. (2006). *Opinion Mining and Sentiment Analysis. Foundations and Trends® in Information Retrieval*, 1(2). 91–231.
- Tata, S., & Di Eugenio, B. (2010). Generating Fine-Grained Reviews of Songs from Album Reviews. *Proceedings of the 48th Annual Meeting of the Association for Computational Linguistics*, (July), 1376–1385.
- Ruihai Dong, Michael P O'Mahony, and Barry Smyth (2014). Further Experiments in Opinionated Product Recommendation. In ICCBR'14, pages 110–124.
- Oramas S., Espinosa-Anke L., Lawlor A., Serra X., Saggion H. (2016). *Exploring Music Reviews for Music Genre Classification and Evolutionary Studies*. 17th International Society for Music Information Retrieval Conference. ISMIR 2016.
- Dominique Moisi. The Geopolitics of Emotion: How Cultures of Fear, Humiliation, and Hope are Reshaping the World. Anchor Books, New York, NY, USA, 2010.

# Semantics

# Knowledge Repositories and Knowledge Bases

- A Knowledge Base (KB) is a rich form of Knowledge Repository (KR), term coined to differentiate from traditional *databases*.
- The term Knowledge Base may be used to refer to lexical databases, ontologies, and any graph-like Knowledge Repository.
- KBs are essential for AI tasks such as reasoning, inference or semantic search. Also for Word Sense Disambiguation, Entity Linking, Semantics...

# Knowledge Bases

- Hand-crafted KBs
  - From generic to domain-specific. E.g. **WordNet**, **CheBi**, **SnomedCT**.
- Integrative Projects
  - Unify in one single resource manually curated KRs and KBs.
    - ⇒ **BabelNet** (originally, WordNet + Wikipedia), **DBPedia**, **Yago**...
- Open Information Extraction for KB construction
  - **NELL**, **PATTY**, **WiseNet**, **DefIE**, **KB-Unify**...

# Music Knowledge Bases

- **MusicBrainz and Discogs**

- Open encyclopedias of music metadata
- MB is regularly published as Linked Data by the LinkedBrainz project.

- **Grove Music Online**

- Music *scholar* encyclopedia

# Information Extraction

Information extraction (IE) is the task of automatically extracting **structured** information from **unstructured** and/or semi-structured machine-readable documents.

Unstructured vs. Structured



# Information Extraction related Tasks

- Entity Linking
- Relation Extraction
- Knowledge Base Construction
- Knowledge Graphs

# Information Extraction

Unstructured text

“Hate It Here” was written by Wilco frontman , Jeff Tweedy .

# Information Extraction

## Entity Identification

“Hate It Here” was written by [Wilco](#) frontman , [Jeff Tweedy](#) .

# Information Extraction

## Entity Recognition



# Information Extraction

## Wilco (disambiguation)

---

From Wikipedia, the free encyclopedia

**Wilco** is an American rock band.

**Wilco** may also refer to:

- [Wilco \(voice procedure\)](#), a radio procedure word, short for "**Will Comply**"; origin of the term
- [Wilco \(The Album\)](#), an album by the band Wilco, or the title song, "Wilco (The Song)"
- [Wilco: Learning How to Die](#), a book about the band, by Greg Kot
- [Wilco \(farm supply cooperative\)](#), an American chain of agricultural cooperative stores
- [Wilco \(tree\)](#), *Anadenanthera colubrina*, a South American tree
- [Wilkinson County, Georgia](#), sometimes abbreviated as "Wilco"
- [Williamson County, Texas](#), sometimes abbreviated as "Wilco"
- [WilcoHess](#), the chain of gas stations

# Information Extraction

Entity Linking or Disambiguation



WIKIPEDIA  
The Free Encyclopedia

Organization

“Hate It Here” was written by Wilco frontman , Jeff Tweedy .

↓  
Work of art

<http://musicbrainz.org/recording/246500ae-379b-4290-8716-d58b596753dd>

↑  
Person

[https://en.wikipedia.org/wiki/Jeff\\_Tweedy](https://en.wikipedia.org/wiki/Jeff_Tweedy)



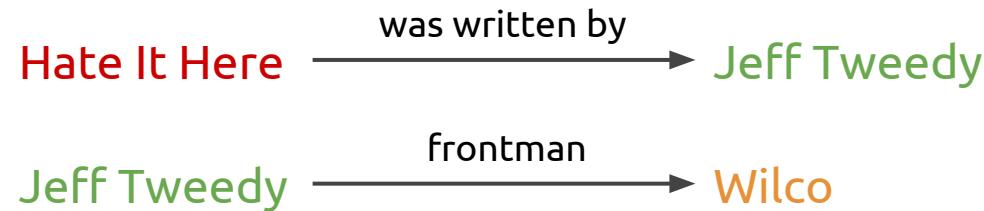
**Hate It Here**  
~ Recording by Wilco



# Information Extraction

## Relation Extraction

“Hate It Here” was written by Wilco frontman , Jeff Tweedy .



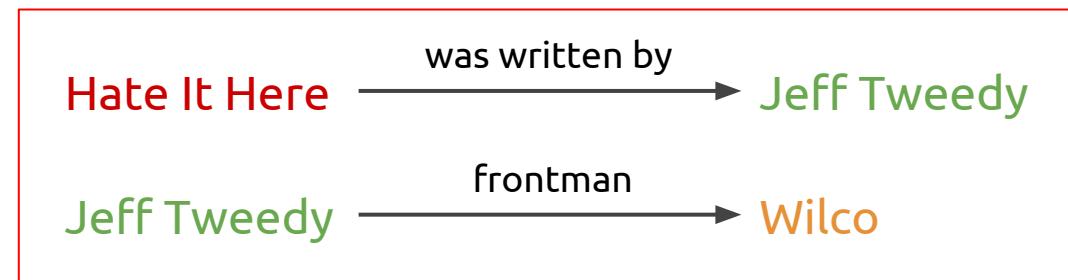
# Information Extraction

## Relation Extraction

Unstructured

“Hate It Here” was written by Wilco frontman , Jeff Tweedy .

Structured



# Entity Linking

Entity linking is the task to associate, for a given candidate textual fragment, the most suitable entry in a reference Knowledge Base.

- Also referred to as Entity Disambiguation
- Typically Wikipedia, DBpedia, YAGO, Freebase as reference KB



# Entity Linking

Entity linking needs to handle:

- **Name variations** (entities are referred to in many different ways)
  - e.g. Elvis, Elvis Presley, Elvis Aaron Presley, The King of Rock and Roll
- **Entity ambiguity** (the same string can refer to more than one entity)
  - e.g. Prince, Debut, Bach, Strauss
- **Missing entities** (there is no target entity in the knowledge base)
  - e.g. Supertrópica is not in Wikipedia

# Entity Linking

Entity linking needs to handle:

- **Name variations** (entities are referred to by different names)
  - e.g. Elvis, Elvis Presley, Elvis Aaron Presley
- **Entity ambiguity** (the same string can refer to multiple entities)
  - e.g. Prince, Debut, Bach, Strauss
- **Missing entities** (there is no target entity in the knowledge base)
  - e.g. Supertrópica is not in Wikipedia



# Entity Linking: Tools

**Babelfy**: Entity Linking + Word Sense Disambiguation. Web service. KB: BabelNet. <http://babelfy.org/index>

**Tagme**: Web service. KB: Wikipedia. <https://tagme.d4science.org/tagme/>

**DBpedia Spotlight**. Installable web service. KB: DBpedia.  
<https://github.com/dbpedia-spotlight/dbpedia-spotlight>

**MEL**: Music Entity Linking service. KB: MusicBrainz . <http://mel.mtq.upf.edu>

# Relation Extraction

Detection and classification of **semantic relations** within a set of **entities**.

Numerous **variants**:

- Supervision: {fully, un, semi, distant}-supervision
- Undefined vs. pre-determined set of relations
- Binary vs. n-ary relations

# Relation Extraction

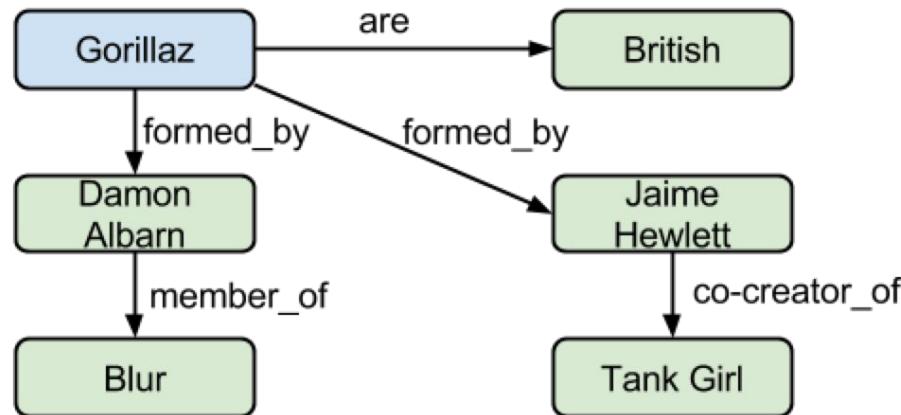
“Gorillaz are a British virtual band formed in 1998 by Damon Albarn of Blur, and Jaime Hawlett, co-creator of the comic book Tank Girl.”

# Relation Extraction

“Gorillaz are a British virtual band formed in 1998 by Damon Albarn of Blur, and Jaime Hawlett, co-creator of the comic book Tank Girl.”

# Relation Extraction

“Gorillaz are a British virtual band formed in 1998 by Damon Albarn of Blur, and Jaime Hewlett, co-creator of the comic book Tank Girl.”



# Relation Extraction (Tools)

**ReVerb**: OpenIE. Downloadable JAR. <http://reverb.cs.washington.edu/>

**OpenIE**: Successor of ReVerb. Downloadable JAR. <http://openie.allenai.org/>

**DeepDive**: Distant supervision. Installable python app.

<http://deepdive.stanford.edu/>

# Relation Extraction (References)

## Traditional IE

Zhao, S., & Grishman, R. (2005). Extracting relations with integrated information using kernel methods. In *Proceedings of the 43rd Annual Meeting on Association for Computational Linguistics - ACL '05* (pp. 419–426).

## Weak Supervision

Bunescu, R. C., & Mooney, R. J. (2007). Learning to Extract Relations from the Web using Minimal Supervision. *Computational Linguistics*, 45(June), 576–583.

## Self Supervision

Carlson, A., Betteridge, J., & Kisiel, B. (2010). Toward an Architecture for Never-Ending Language Learning. In *Proceedings of the Conference on Artificial Intelligence (AAAI) (2010)*

## Distant Supervision

Riedel, S., Yao, L., & McCallum, A. (2010). Modeling relations and their mentions without labeled text. In *Lecture Notes in Computer Science* (Vol. 6323 LNAI, pp. 148–163).

# Relation Extraction (References)

## Open IE

Fader, A., Soderland, S., & Etzioni, O. (2011). Identifying relations for open information extraction. *Proceedings of the Conference on Empirical Methods in Natural Language Processing EMNLP '11*, 1535–1545.

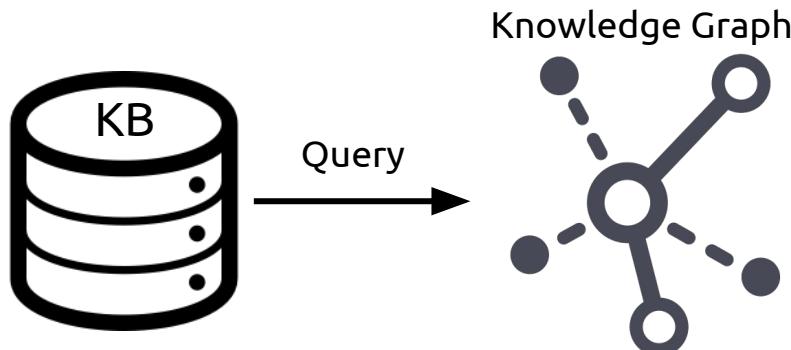
## Semantic Open IE

Nakashole, N., Weikum, G., & Suchanek, F. M. (2012). PATTY: A Taxonomy of Relational Patterns with Semantic Types. *EMNLP-CoNLL*, (July), 1135–1145.

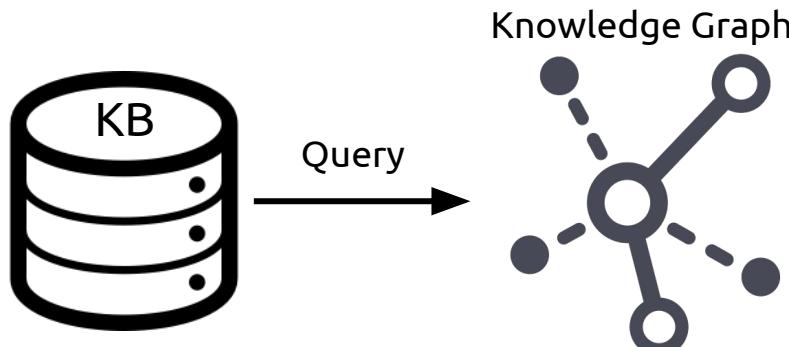
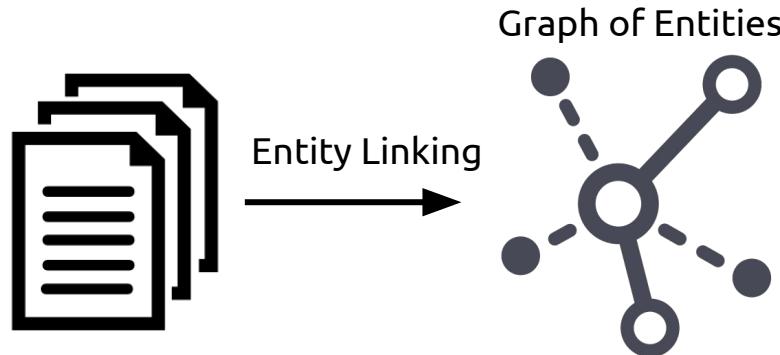
Delli Bovi, C., Telesca, L., & Navigli, R. (2015). Large-Scale Information Extraction from Textual Definitions through Deep Syntactic and Semantic Analysis. *Transactions of the Association for Computational Linguistics*, 3, 529–543.

Oramas S., Espinosa-Anke L., Sordo M., Saggion H., Serra X. Information Extraction for Knowledge Base Construction in the Music Domain. *Journal on Knowledge & Data Engineering*, Elsevier.

# Knowledge Graphs



# Knowledge Graphs



# Graph of Entities

**Gorillaz BIO:** “[Gorillaz](#) are a [British](#) virtual band formed in 1998 by [Damon Albarn](#) of [Blur](#), and [Jaime Hawlett](#), co-creator of the comic book [Tank Girl](#).”

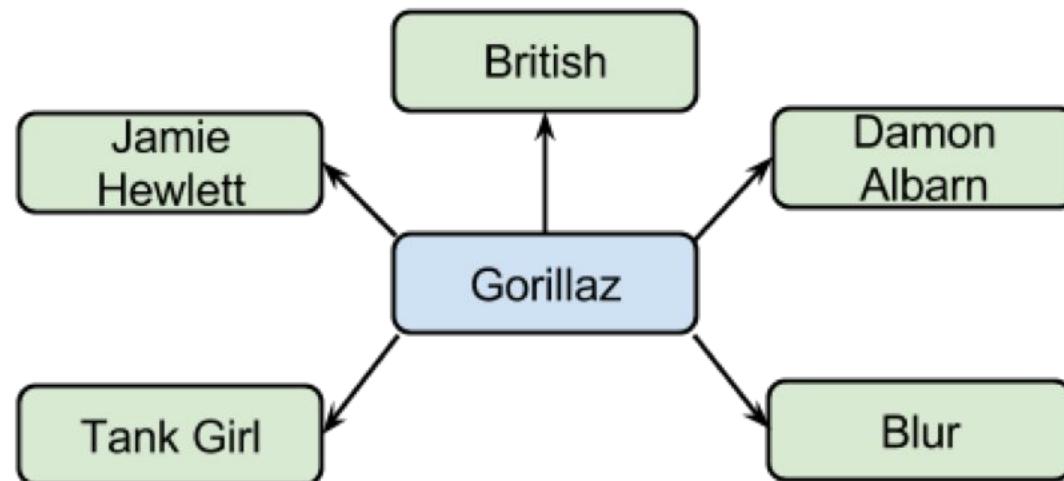
# Graph of Entities

**Gorillaz BIO:** “[Gorillaz](#) are a [British](#) virtual band formed in 1998 by [Damon Albarn](#) of [Blur](#), and [Jaime Hawlett](#), co-creator of the comic book [Tank Girl](#).”

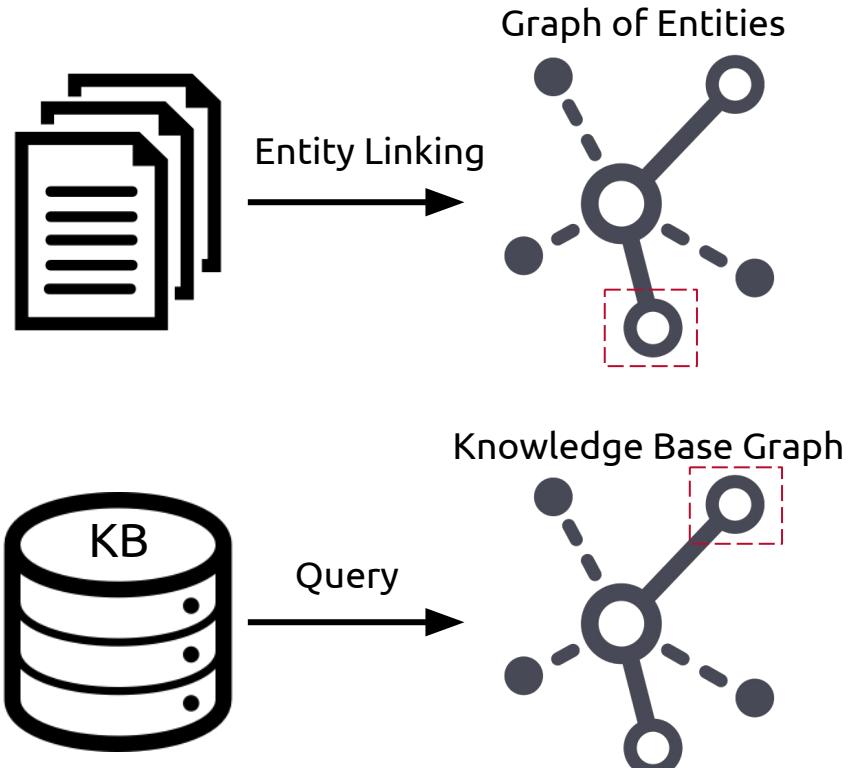
Gorillaz

# Graph of Entities

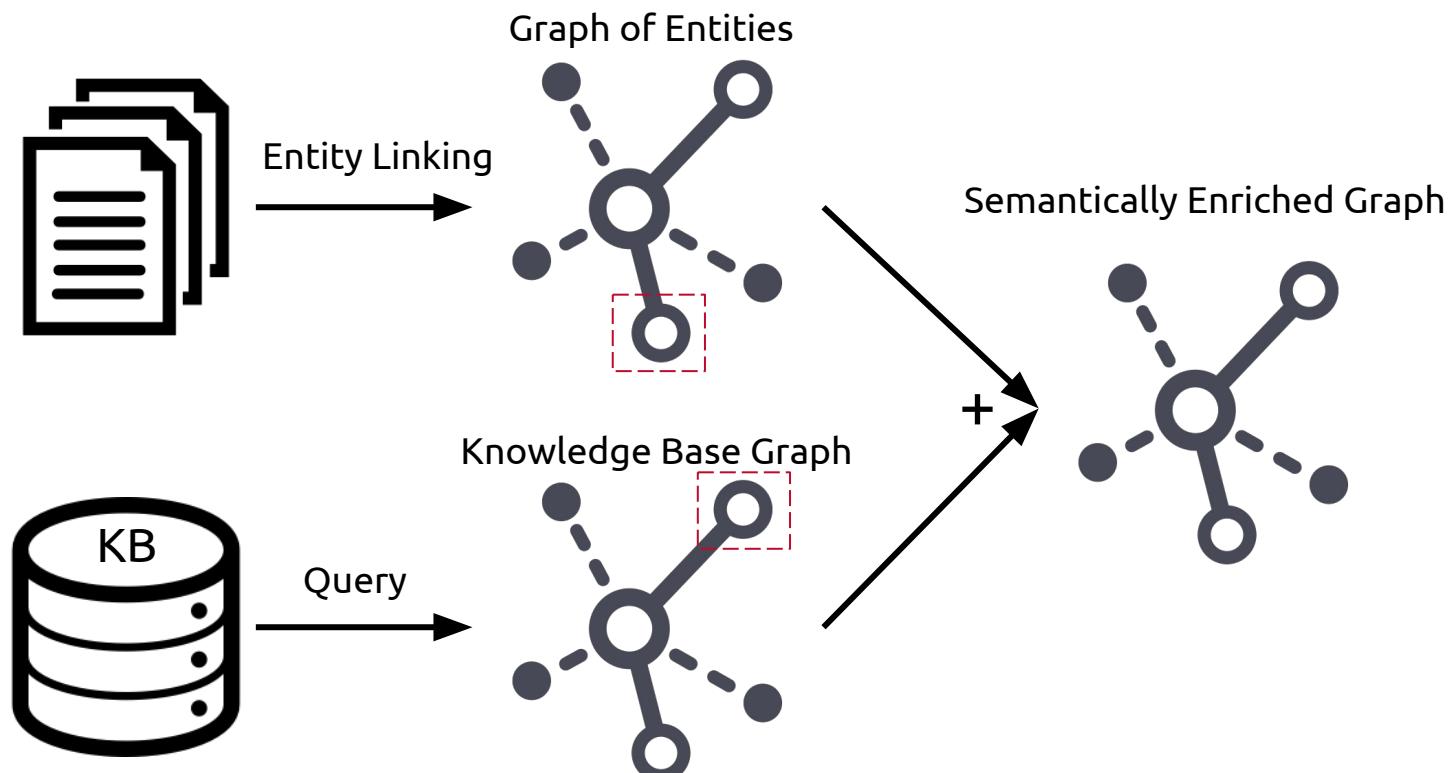
**Gorillaz BIO:** “[Gorillaz](#) are a [British](#) virtual band formed in 1998 by [Damon Albarn](#) of [Blur](#), and [Jaime Hewlett](#), co-creator of the comic book [Tank Girl](#).”



# Knowledge Graphs

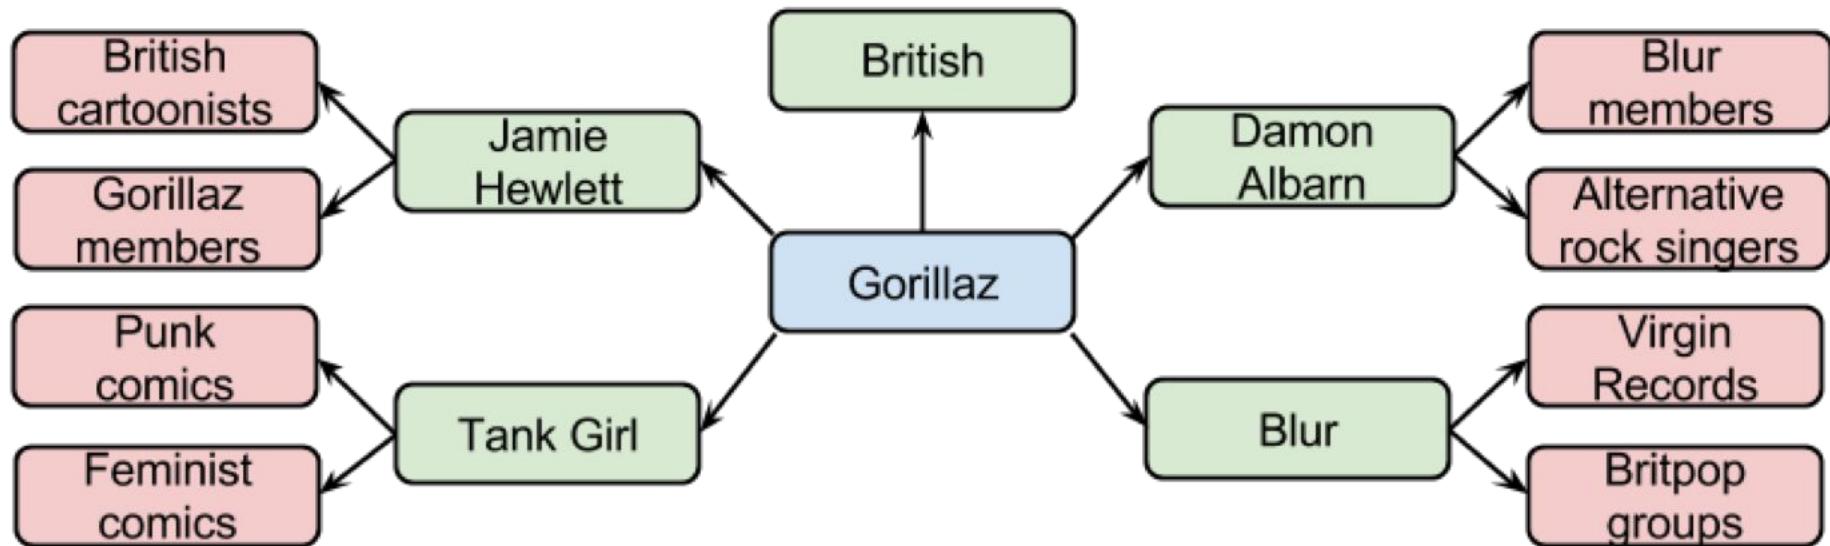


# Knowledge Graphs



# Semantically Enriched Graph

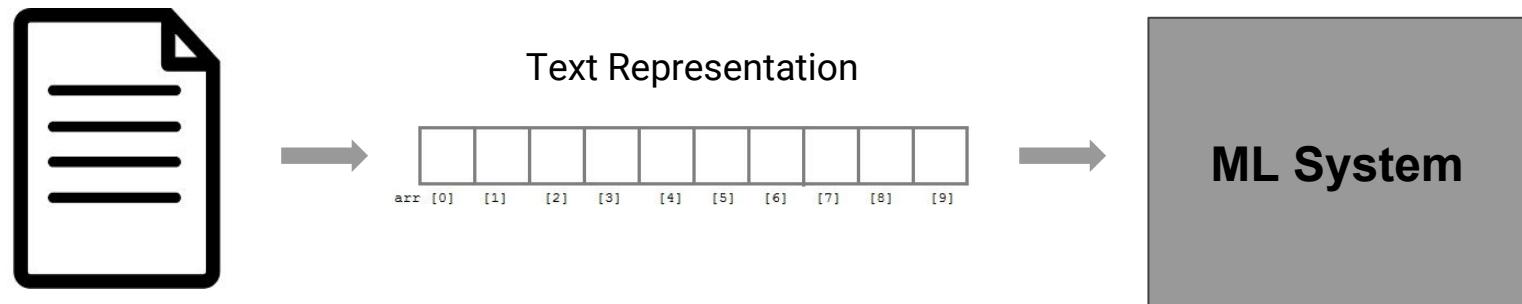
**Gorillaz BIO:** “[Gorillaz](#) are a [British](#) virtual band formed in 1998 by [Damon Albarn](#) of [Blur](#), and [Jaime Hewlett](#), co-creator of the comic book [Tank Girl](#).”



# Text Representation

# Text Representation

To process a text in a Machine Learning system we need to create a suitable data representation, which is typically a vector of real numbers.



# Text Representation

Types of representation:

- Vector Space Model (VSM)
- Graph embeddings
- Word embeddings
- Char embeddings

# Vector Space Model (VSM)

**Bag-of-words (BoW):** unordered list of words

Red Hot Chilli Peppers played at the red club →



# Vector Space Model (VSM)

**VSM:** Given the BoW extracted from the document, create a feature vector where each feature represents a word and the feature value is a term weight.

The term weight might be:

- a binary value (1 if word occurred in the document, and 0 if did not)
- a term frequency value (how many times the word occurred in the document)
- a numerical statistic value (e.g., tf-idf, BM25)

# tf-idf

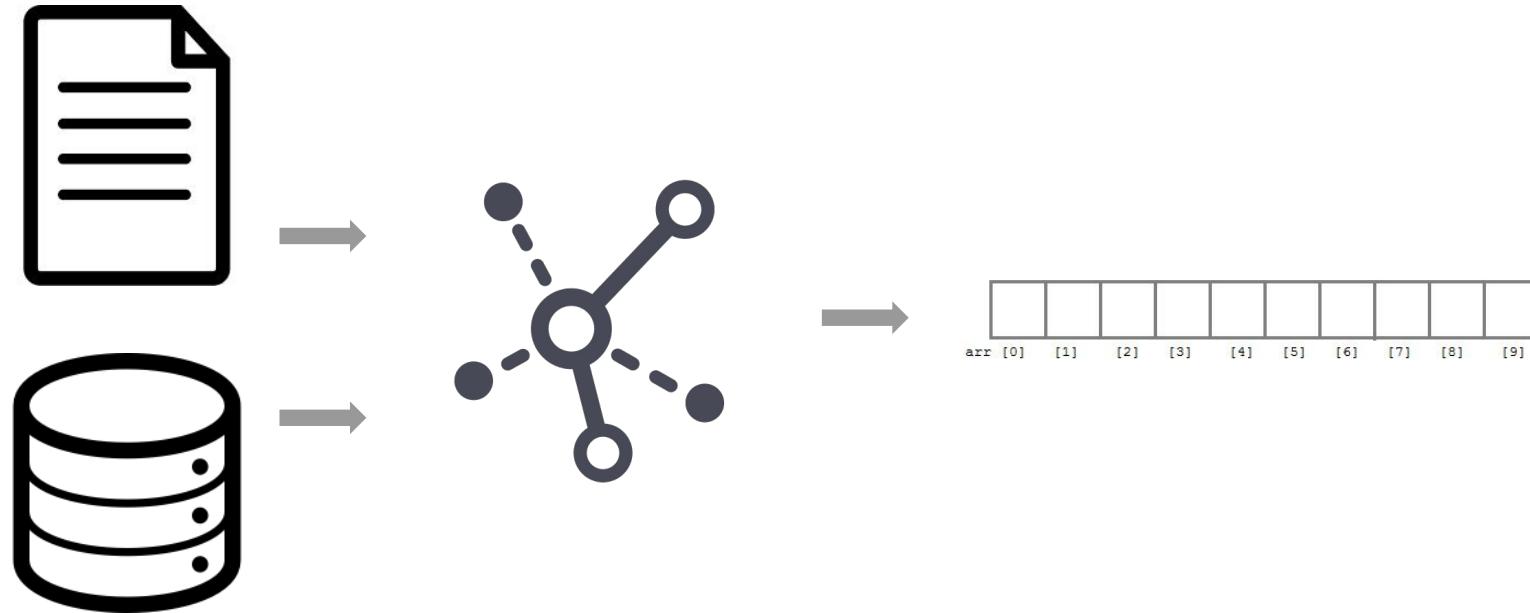
## tf-idf: term frequency-inverse document frequency

- Numerical statistic that is intended to reflect how important a word is to a document in a collection or corpus
- Widely used for Information Retrieval (search) and text-mining
- Frequently found in MIR literature

$TF(t) = (\text{Number of times term } t \text{ appears in a document}) / (\text{Total number of terms in the document})$

$IDF(t) = \log_e(\text{Total number of documents} / \text{Number of documents with term } t \text{ in it})$

# Graph Embedding



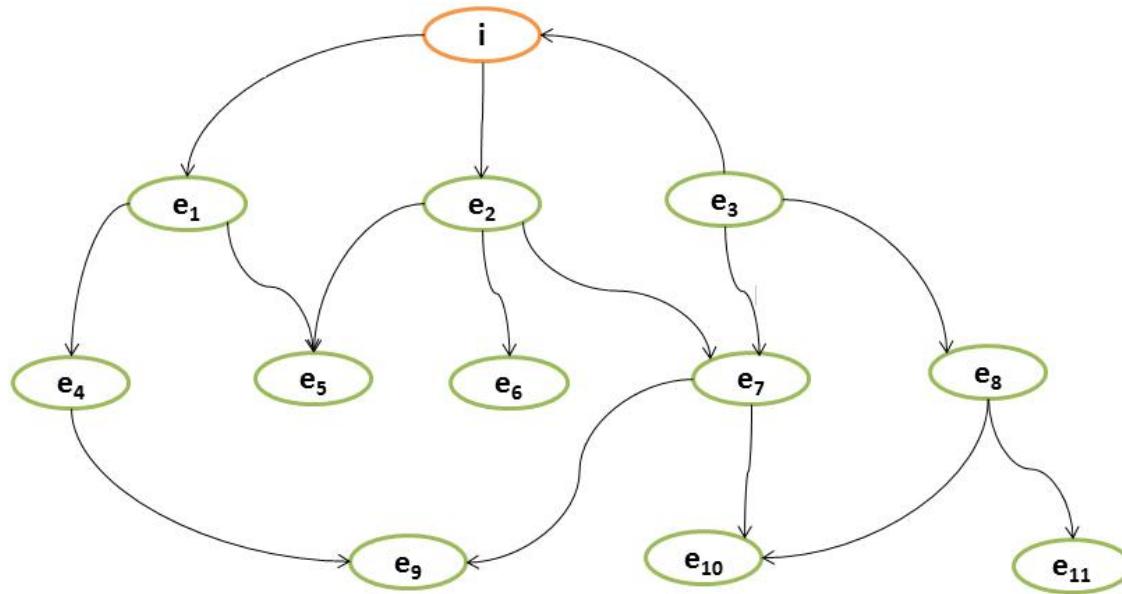
# Graph Embeddings

Represent in a linear vector the structure of the graph

Multiple approaches:

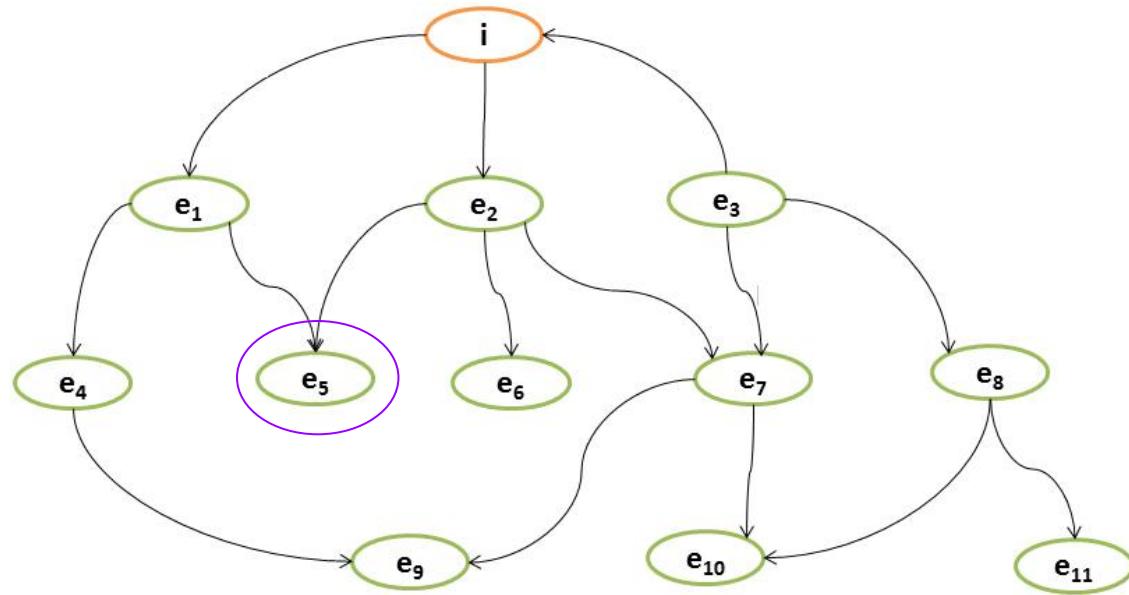
- Node and path embeddings
- Adjacency matrix
- Statistical models (e.g., node2vec)

# Node-based Embedding



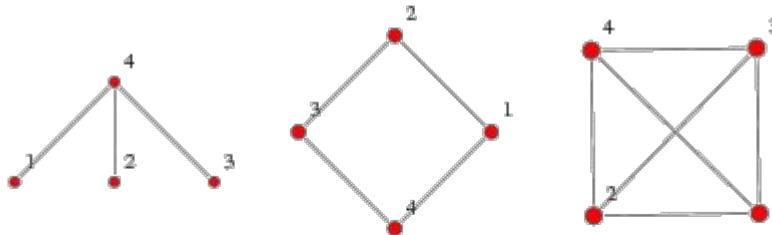
- One **feature** per entity
- **Weight** according to:
  - Distance to root
  - Number of in-links

# Node-based Embedding



- One **feature** per entity
- **Weight** according to:
  - Distance to root
  - Number of in-links

# Adjacency Matrix



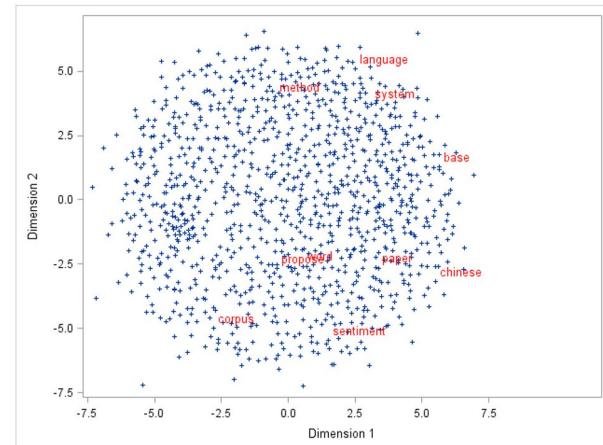
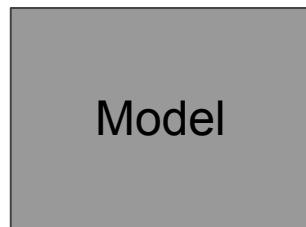
$$\begin{pmatrix} 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix}$$

$$\begin{pmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{pmatrix}$$

$$\begin{pmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix}$$

# Statistical models

Each node becomes a point in an N-dimensional space (node2vec)



# Word Embeddings

## Distributional Semantics

“You shall know a word by the company it keeps”, Firth (1957).

**Distributional Hypothesis:** words that appear in similar contexts exhibit similar semantics.

He filled the **wampimuk** with the substance, passed it around we all drunk some.

We found a little, hairy **wampimuk** sleeping behind the tree.

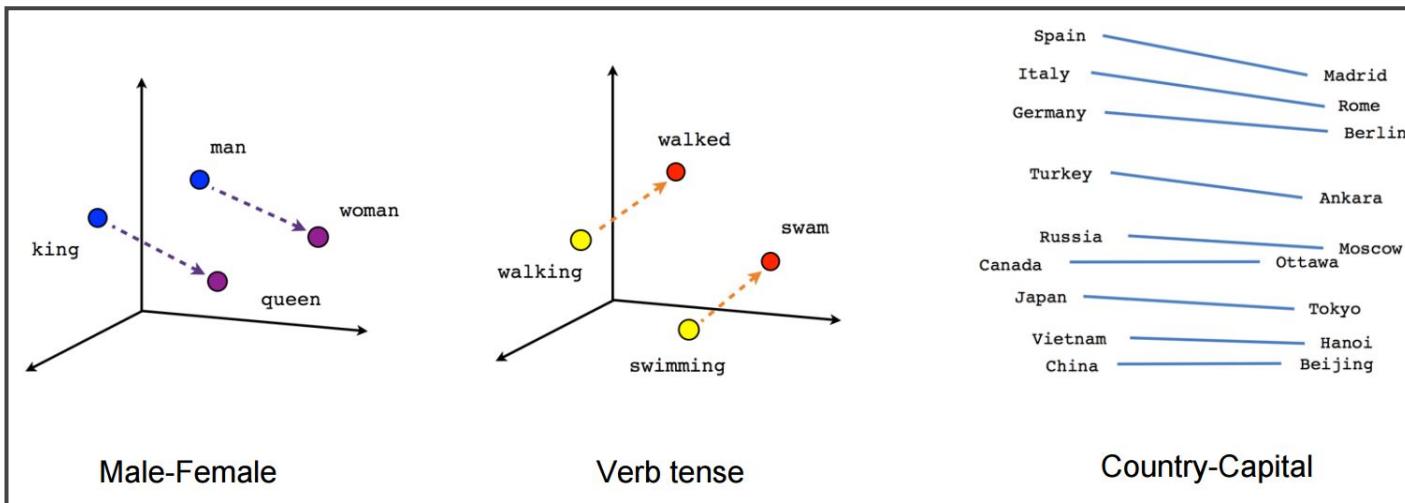
(McDonald and Ramscar, 2001)

# Word Embeddings

## Properties

**Objective:** Project linguistic items into a vector space

Similarity and geometrical properties



# Word Embeddings

## Word Embeddings

Pre-trained on external large corpora

Used as features in different tasks and domains

Popular approaches

word2vec ([Mikolov et al., 2013](#))

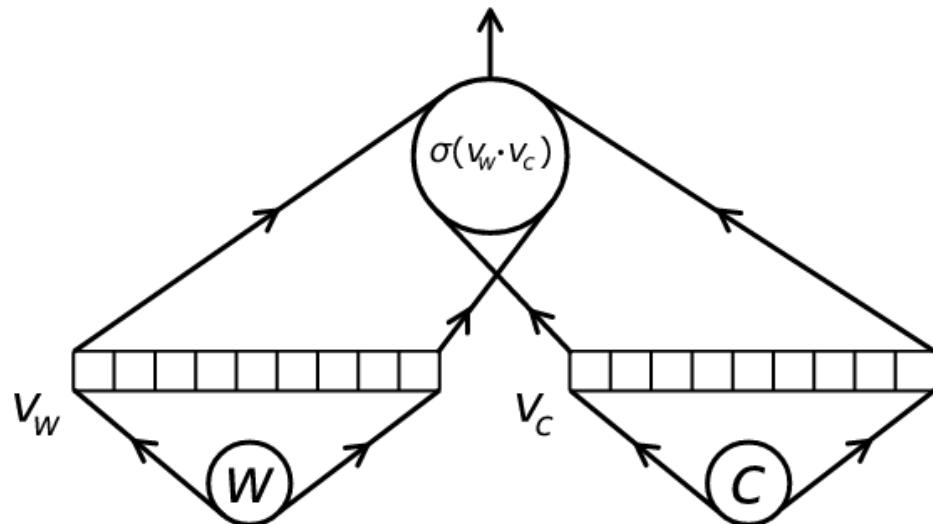
Glove ([Pennington et al., 2014](#))

fastText ([Bojanowski et al., 2017](#))

# Word2vec

Predict a context word  $c \in (w_{i-L}, \dots, w_{i-1}, w_{i+1}, \dots, w_{i+L})$  given a word  $w_i$

$$P(D = 1 | w, c) = \sigma(v_w \cdot v_c)$$



$$\sigma(a) = \frac{1}{1+e^{-a}}$$

# More about Word2vec

<https://www.tensorflow.org/versions/r0.10/tutorials/word2vec/index.html>

<http://cgi.cs.mcgill.ca/~enewel3/posts/implementing-word2vec/>

<http://hduongtrong.github.io/2015/11/20/word2vec/>

T. Mikolov et al (2013): Distributed Representations of Words and Phrases and their Compositionality. Advances in neural information processing systems.

O. Levy, Y. Goldberg (2014): Neural Word Embedding as Implicit Matrix Factorization. NIPS 2014

# Distributional Lexical Semantics

```
>>> from gensim.models import Word2Vec
```

```
>>> model = Word2Vec.load(PATH)
```



<https://radimrehurek.com/gensim/models/word2vec.html>

# Word2vec in Playlists

Trained with Gensim in Art of the Mix playlists

(<http://labrosa.ee.columbia.edu/projects/musicsim/aotm.htm>)

```
model.most_similar('miles davis')
[('john cloтрane', 0.88384414), ('dizzie gillespie', 0.78484219), ('charlie walker', 0.74520659)]
```

```
model.most_similar('marilyn manson')
[('godsmack', 0.93274206), ('white zombie', 0.91064525), ('drowning pool', 0.90275443)]
```

```
model.most_similar('nirvana')
[('soundgarden', 0.84231329), ('pearl jame', 0.8271907), ('oysterhead', 0.81855756)]
```

# Character Embeddings

Each character has its own embedding

Word embeddings are combinations of character embeddings

Useful for out-of-vocabulary words

# Deep Learning Architectures

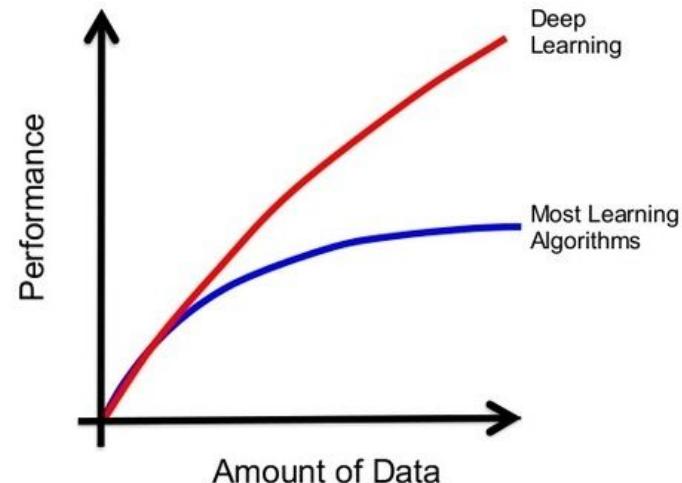
# Deep Learning in Natural Language Processing

Deep Learning improves almost all tasks in NLP!! (as in many other fields)

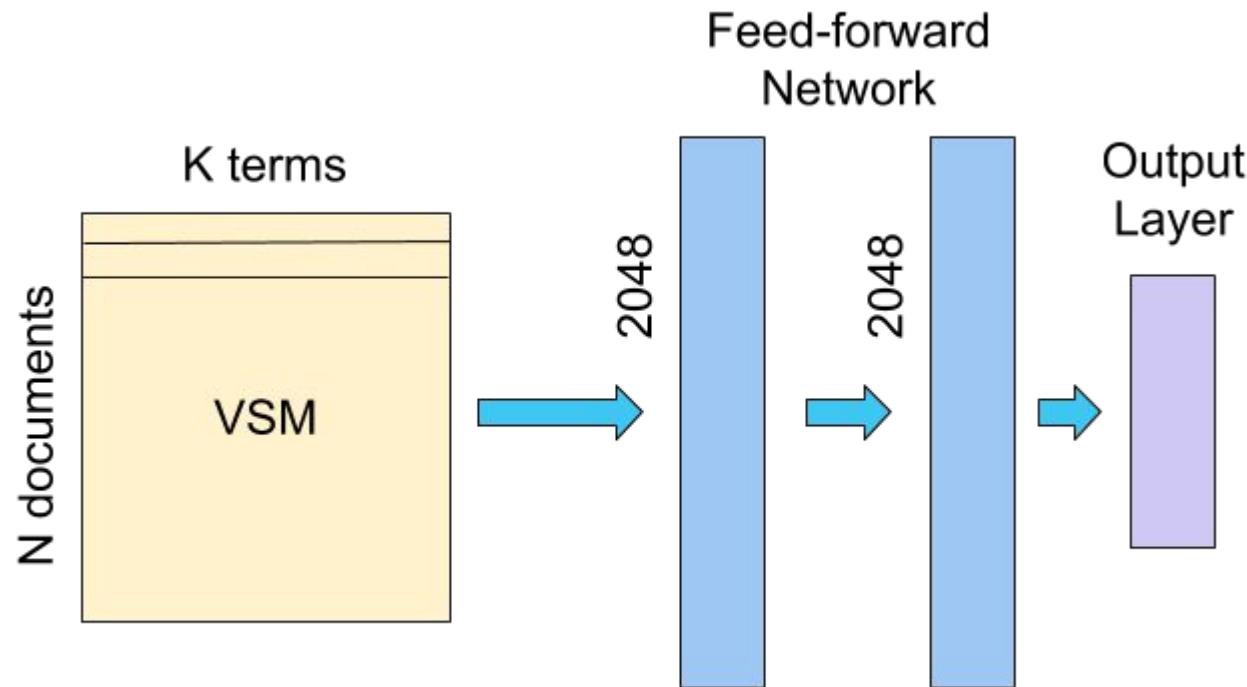
Deep Network Architectures: LSTM y CNN

More than words: character embeddings,  
word embeddings, transfer learning

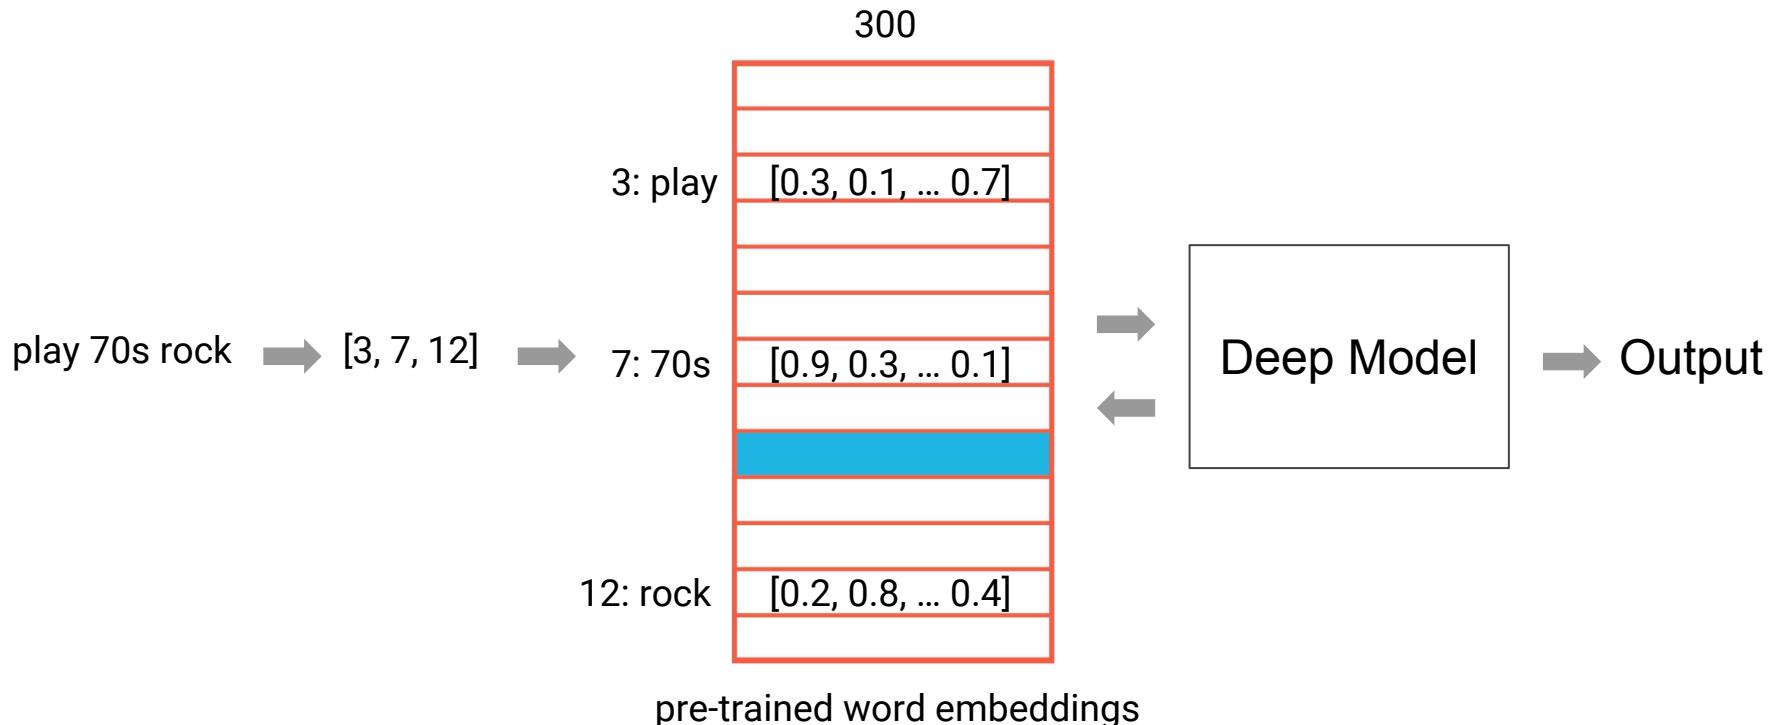
## BIG DATA & DEEP LEARNING



# Feedforward Networks

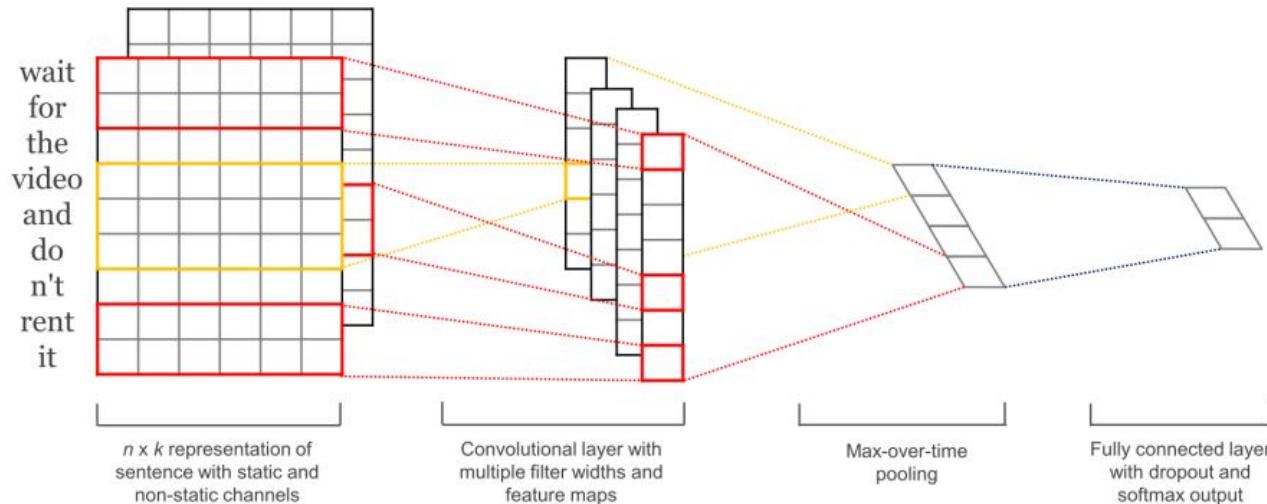


# Embedding Layer

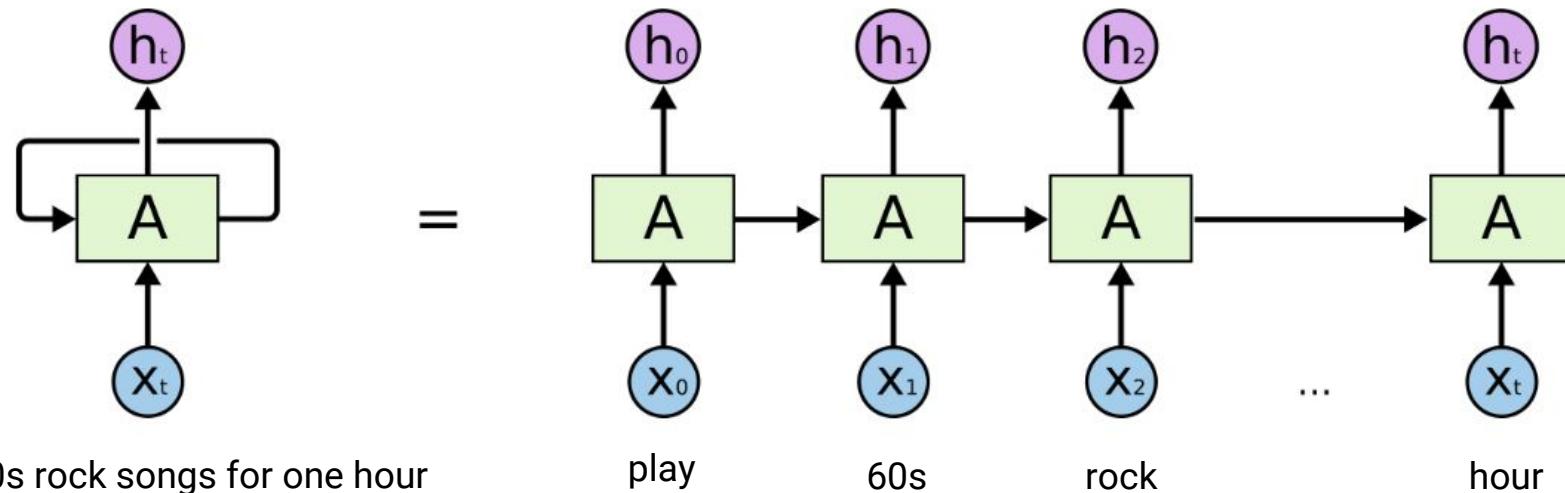


# CNNs

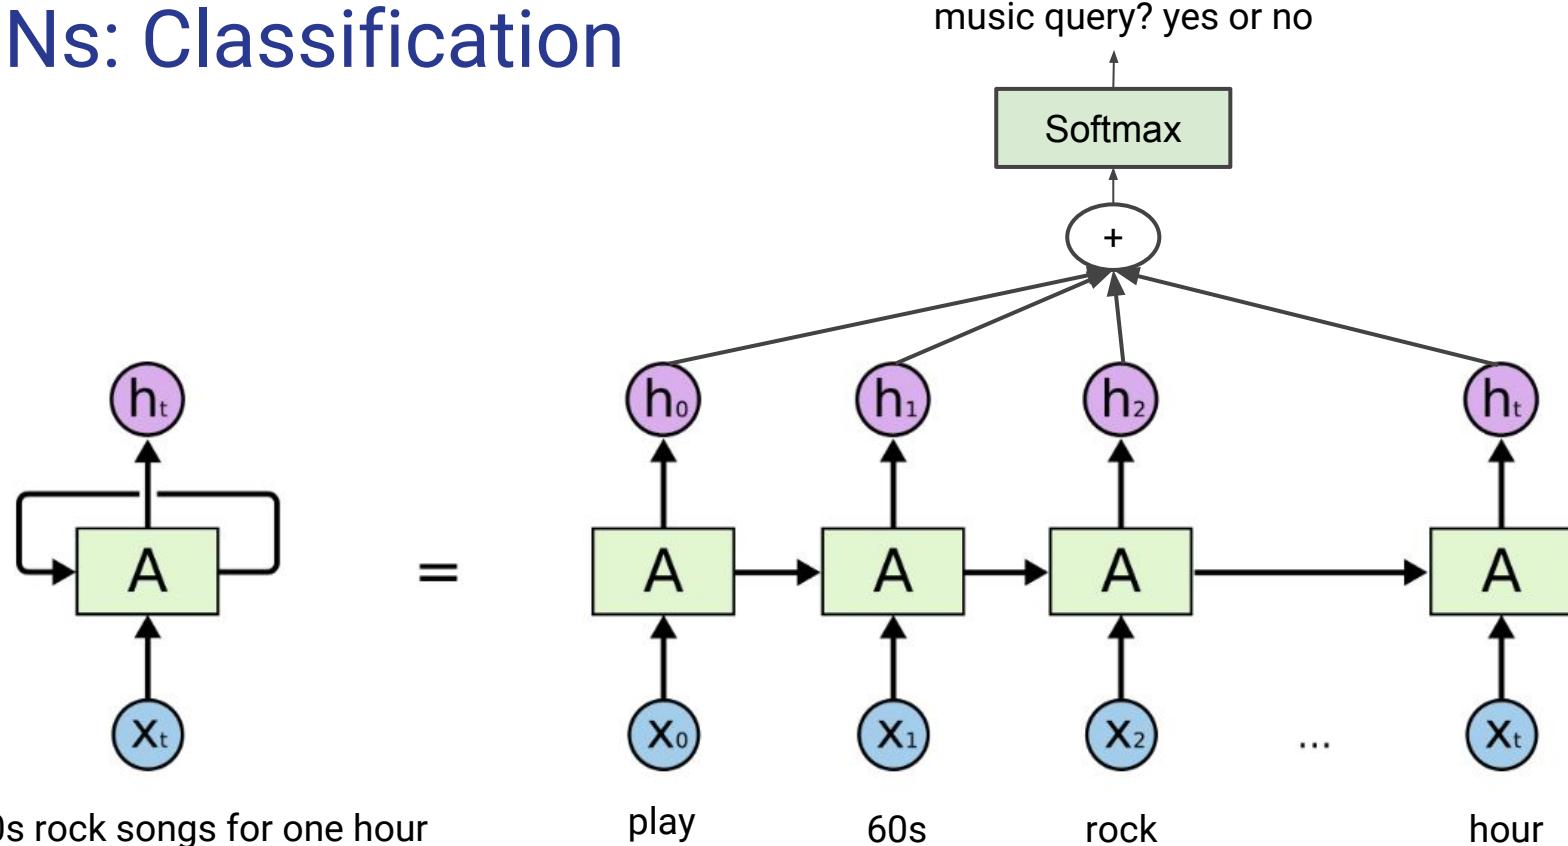
Kim, Y. (2014). Convolutional neural networks for sentence classification. *arXiv preprint arXiv:1408.5882*.



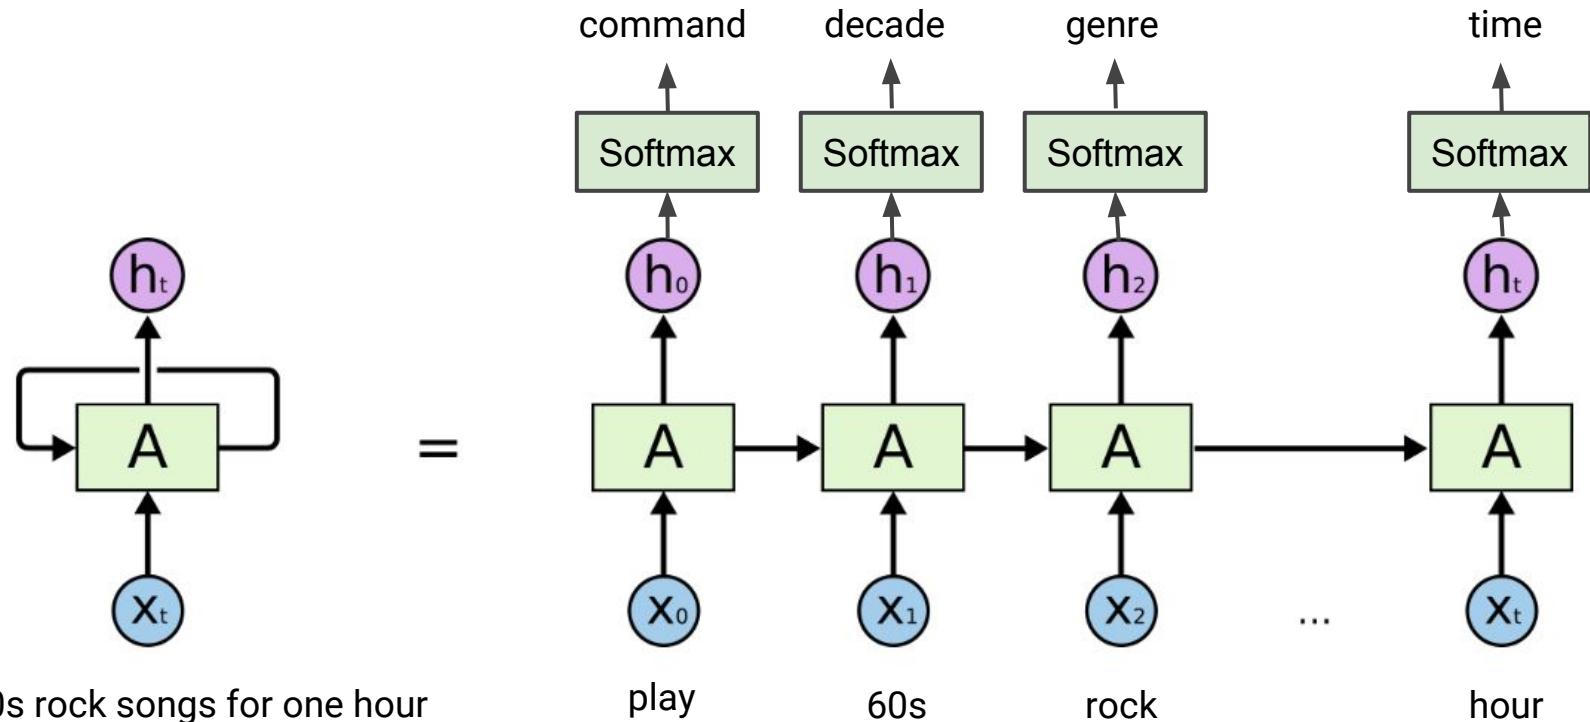
# RNNs (LSTMs, GRUs, vanilla RNNs)



# RNNs: Classification



# RNNs: Seq2seq



Play 60s rock songs for one hour

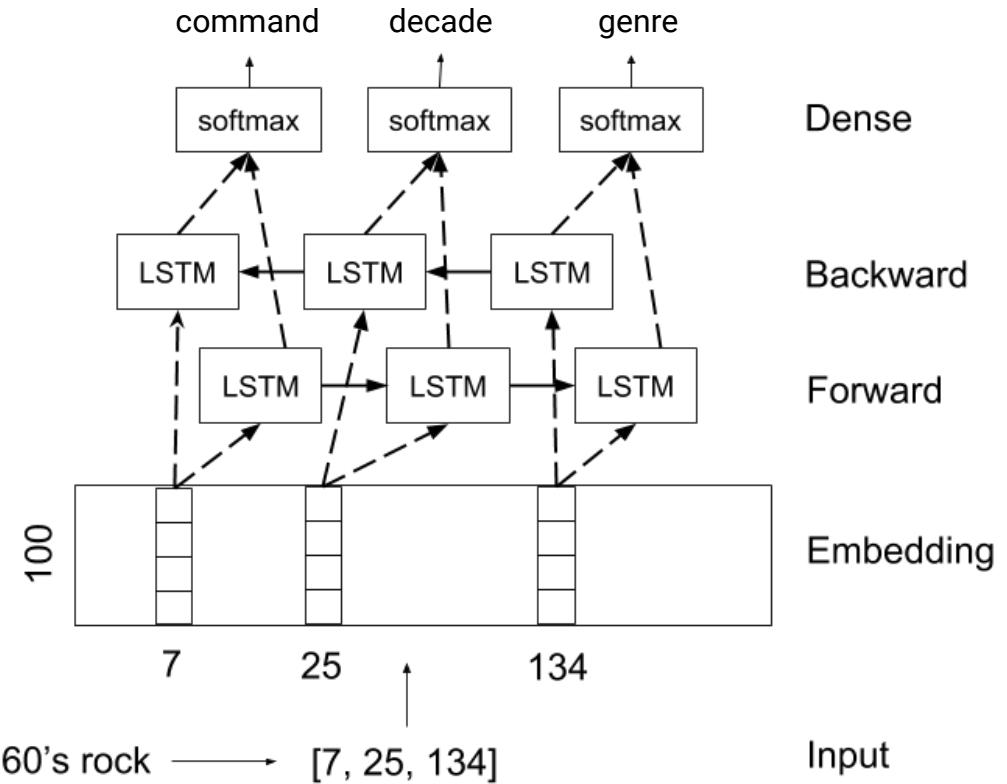
play

60s

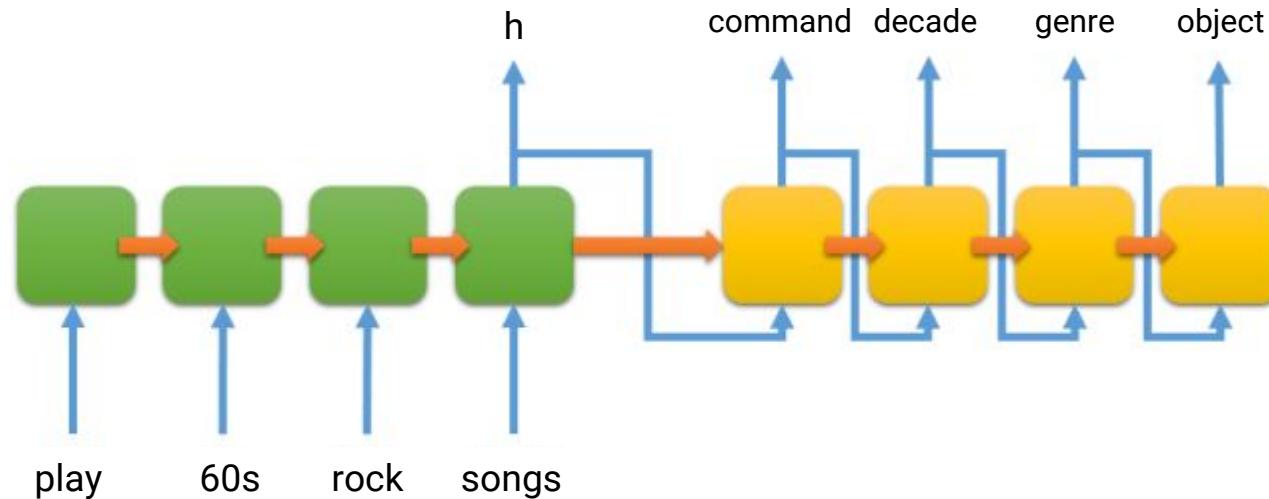
rock

hour

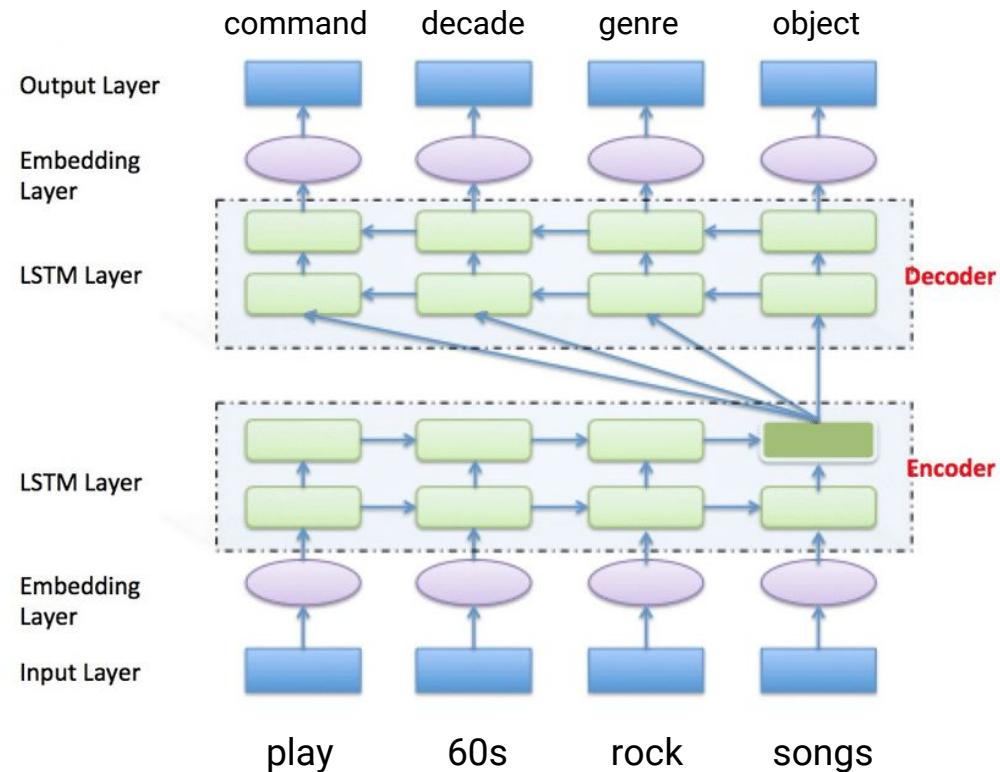
# BiLSTM



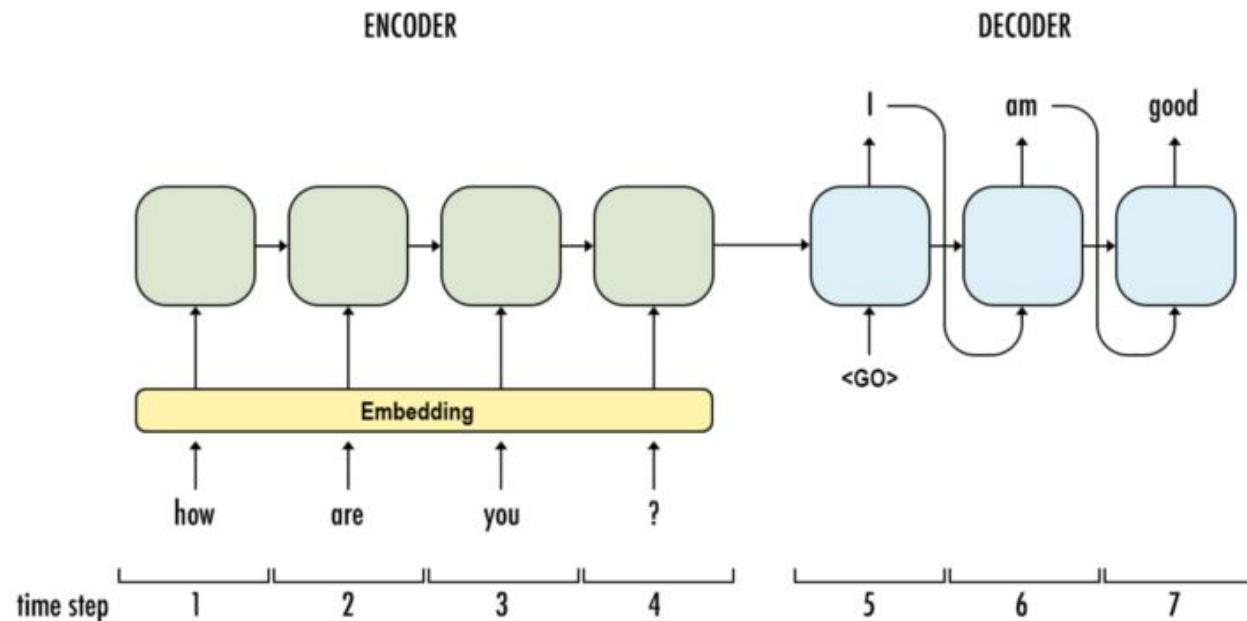
# Encoder-Decoder



# Encoder-Decoder

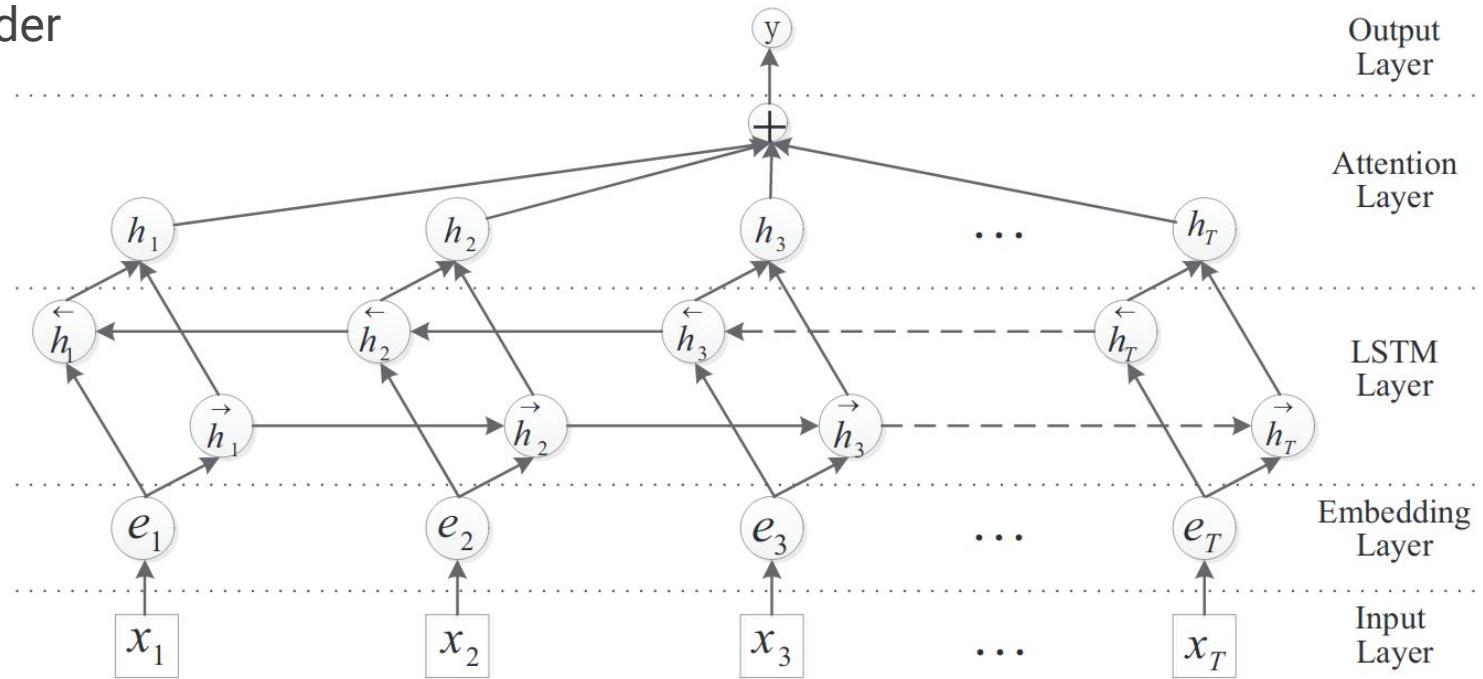


# Encoder-Decoder



# Attention

Encoder

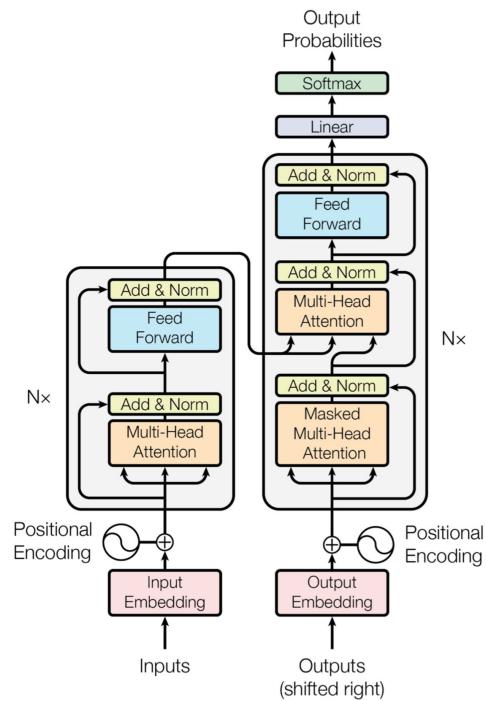


# Transformer

Attention is all you need ([Ashish Vaswani et al. 2017](#))

State-of-the-art in Machine Translation

Improvement on several NLP tasks over RNNs

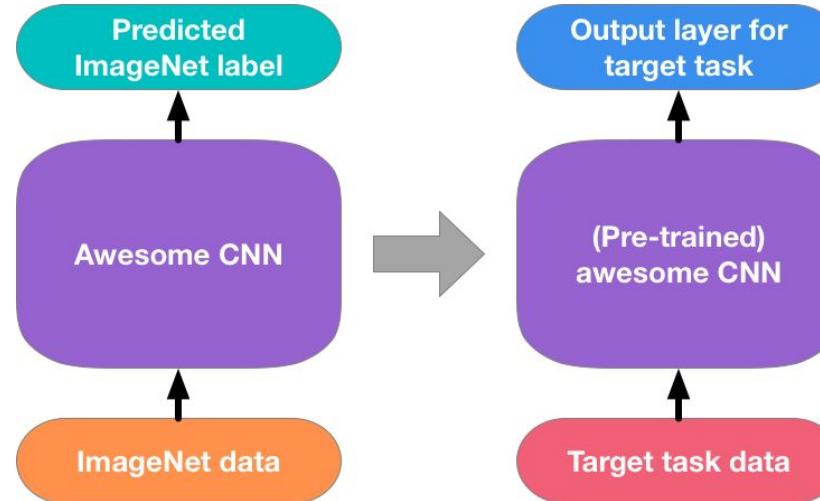


# Contextualized Word Embeddings

# Preliminary Concepts

## Fine-tuning

**Fine-tuning** is the process in which parameters of a model must be adjusted very precisely in order to fit with certain observations.



# Preliminary Concepts

## Language Model

**Language Model:** Probabilistic model that is able to predict the next word in the sequence given the words that precede it.



Almost infinite training corpora

Typically trained using RNNs

Can be used as a feature extractor

Can go forward or backward

# Word Embeddings

## Types

Word Embeddings (one embedding per word)

- ▶ I'm a huge metal fan
- ▶ I went to a heavy metal concert
- ▶ This table is made of metal

Sense Embeddings (one embedding per sense)

- ▶ I'm a huge metal fan
- ▶ I went to a heavy metal concert
- ▶ This table is made of metal

Contextualized Word Embeddings (one embedding per word and context)

- ▶ I'm a huge metal fan
- ▶ I went to a heavy metal concert
- ▶ This table is made of metal

# Word Embeddings

## Word Embeddings

Single vector per word:

First transfer learning approach in NLP,

Pre-trained on external large corpora

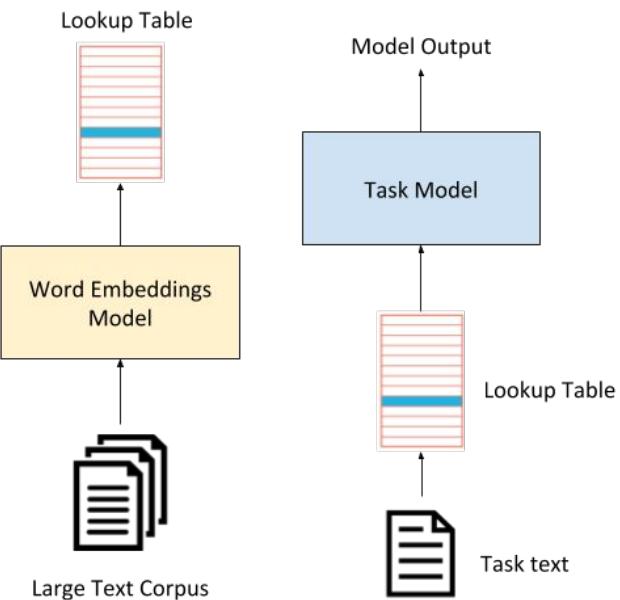
Used as features in different tasks and domains

Popular approaches

word2vec ([Mikolov et al., 2013](#))

Glove ([Pennington et al., 2014](#))

fastText ([Bojanowski et al., 2017](#))



# Word Embeddings

## Sense Embeddings

metal

Several vectors per word (one vector per sense):

Popular approaches:

Nasari: Wikipedia entity embeddings

([José Camacho-Collados et al. 2016](#))

DeConf: WordNet sense embeddings

([M. T. Pilehvar et al. 2016](#))

...



# Word Embeddings

## Contextualized embeddings

Several vectors per word (one vector per context):

Breakthrough in 2018! “ImageNet-like” momentum for NLP

Feature extractor model instead of a lookup table

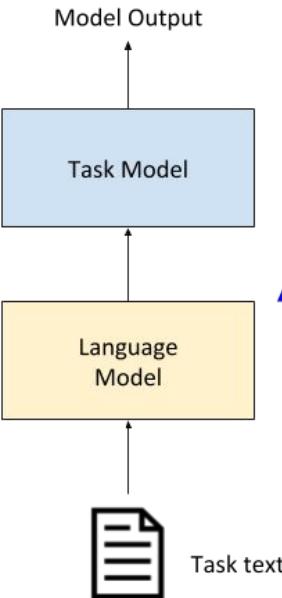
Popular contextualized embedding approaches:

ELMo

ULMFiT

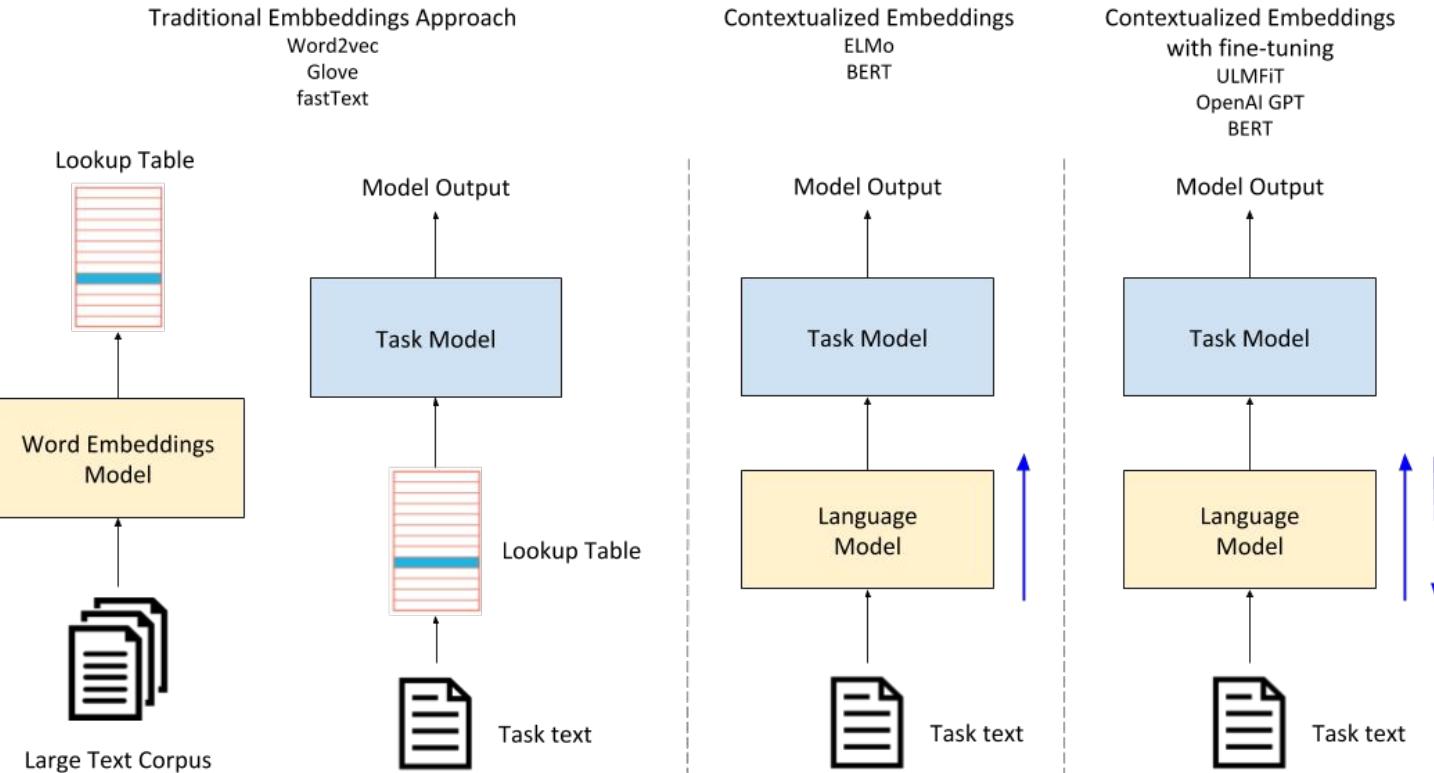
OpenAI GPT

BERT



# Word Embeddings

## Summary



# Further Reading

Interesting blog posts

<https://jalammar.github.io/illustrated-bert/>

<http://jalammar.github.io/illustrated-transformer/>

<https://towardsdatascience.com/beyond-word-embeddings-part-2-word-vectors-nlp-modeling-from-bow-to-bert-4ebd4711d0ec>

# Multimodality

# Multimodality

In large music collections information is multimodal:

audio, images, texts, videos

How can we combine different modalities effectively?



## Biography

**Wilco** is a rock band which formed in 1994 in Chicago, Illinois, United States by the remaining members of alternative country group Uncle Tupelo following singer Jay Farrar's departure from that group. Wilco's lineup has changed frequently, with only singer Jeff Tweedy and bassist John Stiratt remaining from the original incarnation. The other current members are guitarist Nels Cline, multi-instrumentalists Pat Sansone and Mikael Jorgenson, and drummer Glenn Kotche. Wilco has released nine studio albums, a live double album, and three collaborations: two with Billy Bragg, and one with The Minus 5.

**YEARS ACTIVE**  
1994 – present (23 years)

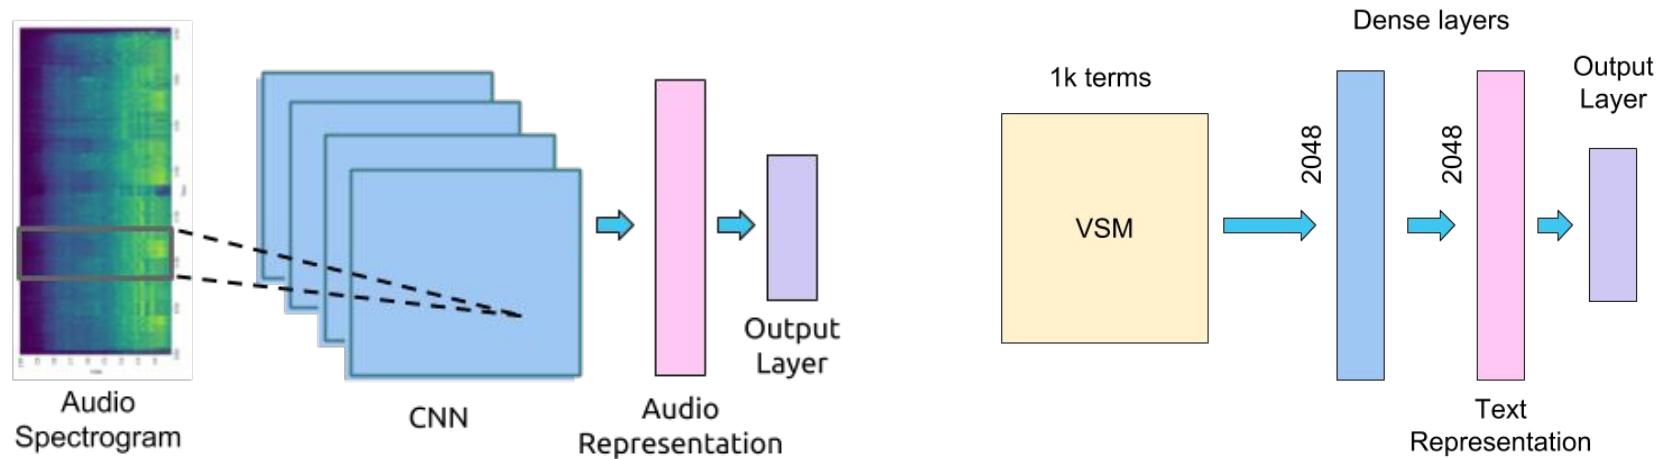
**FOUNDED IN**  
Chicago, Cook County, Illinois, United States

**MEMBERS**  
Bob Egan  
Glenn Kotche (2000 – present)  
Jay Bennett (1995 – 2001)  
Jeff Tweedy



# Learn Data Representations

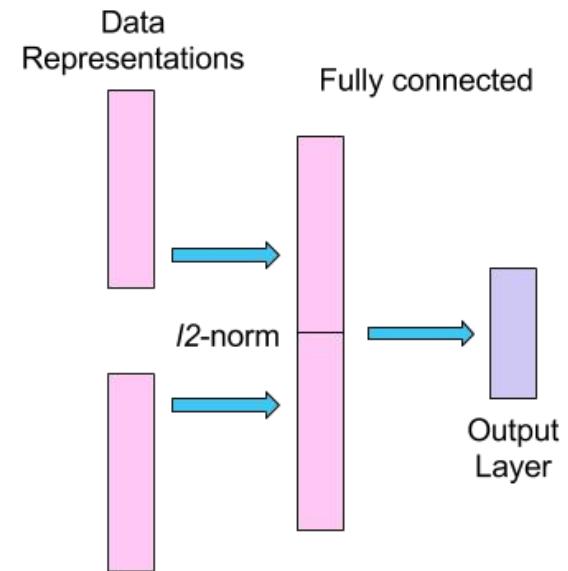
Train on the same task



# Fusion of Data Representations

Extract internal representation from each modality

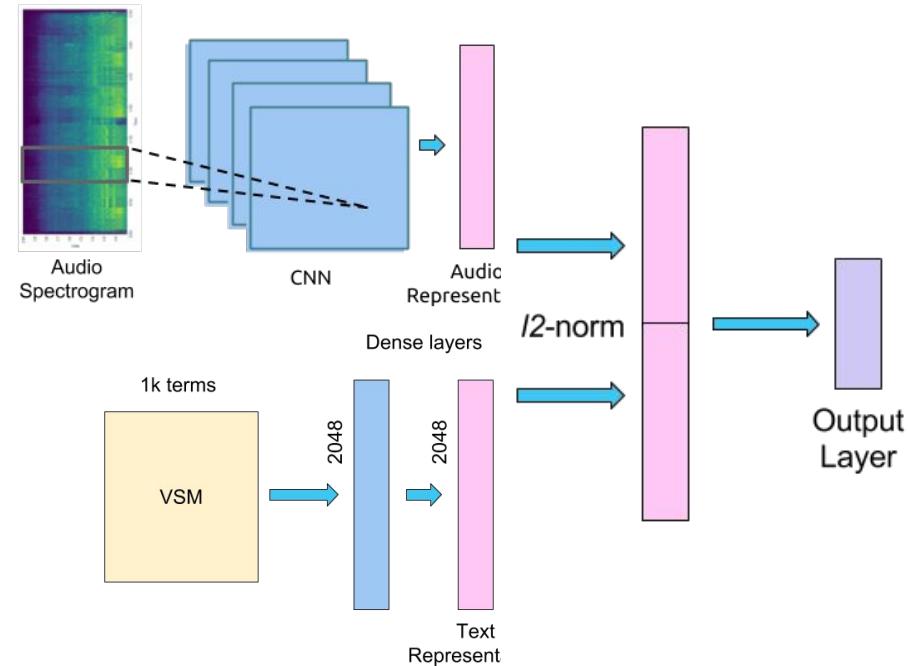
Concatenate and retrain on the same task



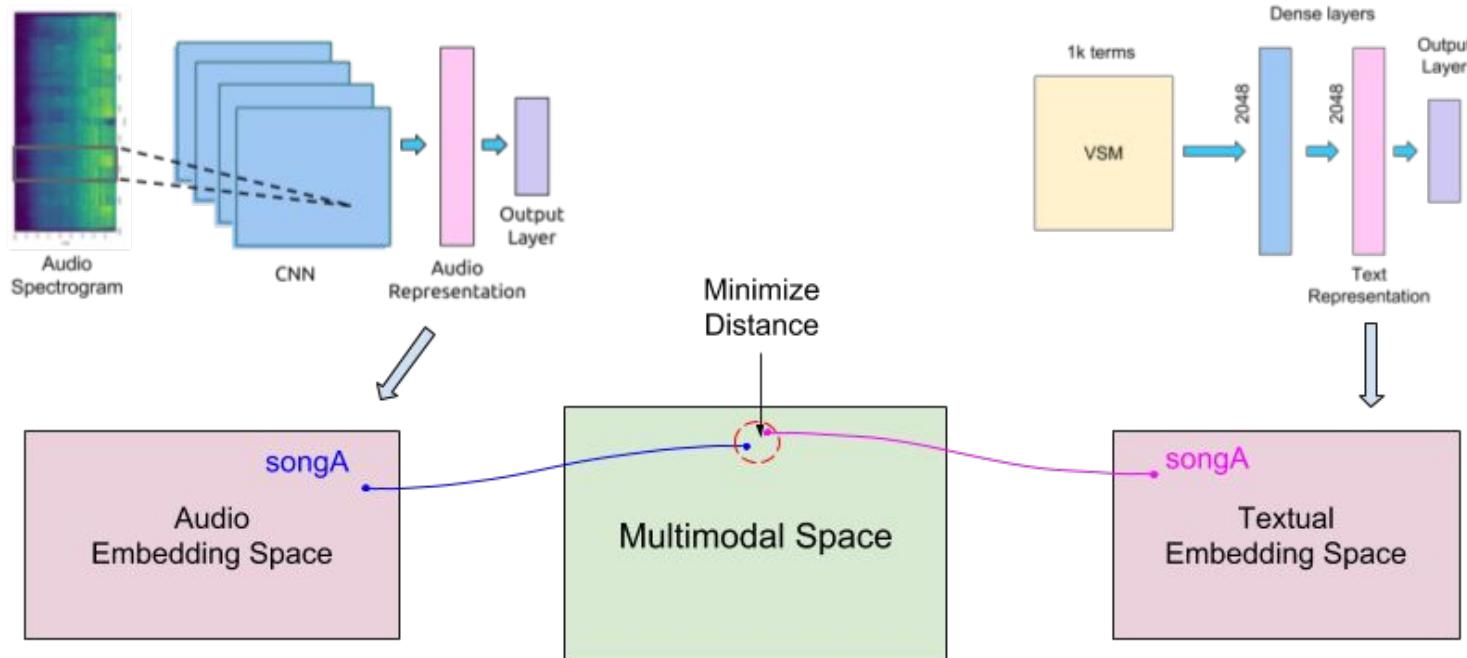
# Fine-tuning of Multimodal Network

Pre-train modalities in isolation

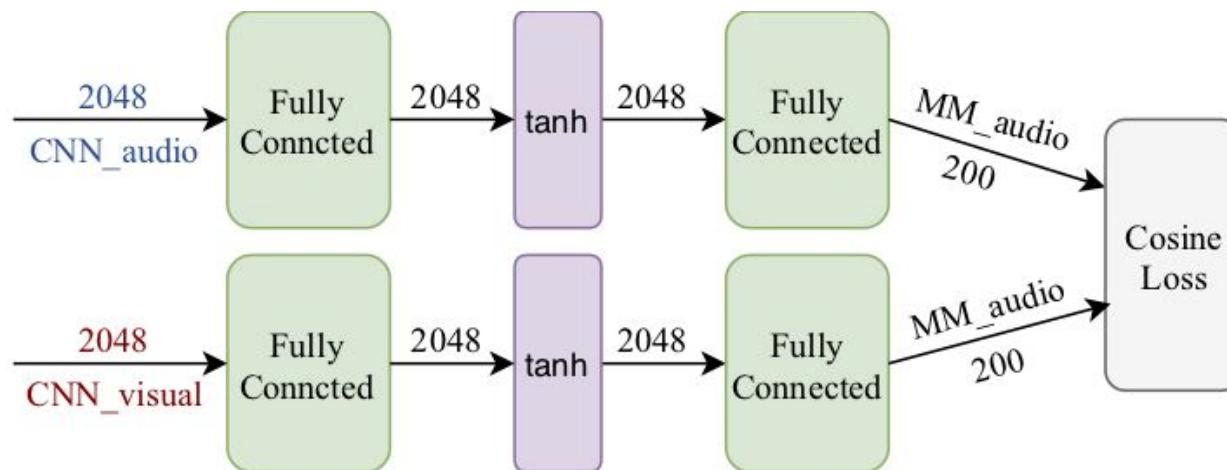
Multimodal Fusion with full networks



# Multimodal Space



# Multimodal Space



# References

Oramas, S., Barbieri F., Nieto O., & Serra X. (2018). Multimodal Deep Learning for Music Genre Classification. *Transactions of the International Society for Music Information Retrieval*. 1(1), 4-21.

Oramas, S., Bogdanov D., & Porter A. (2018). MediaEval 2018 AcousticBrainz Genre Task: A baseline combining deep feature embeddings across datasets. *MediaEval 2018 Workshop*

Oramas, S., Nieto O., Sordo M., & Serra X. (2017). A Deep Multimodal Approach for Cold-start Music Recommendation. *2nd Workshop on Deep Learning for Recommender Systems*, at RecSys 2017

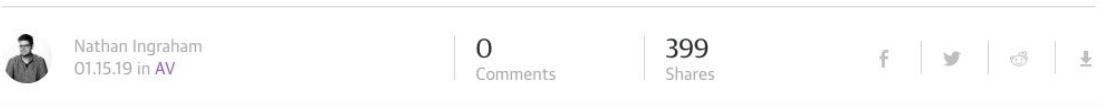
Oramas, S., Nieto O., Barbieri F., & Serra X. (2017). Multi-label Music Genre Classification from Audio, Text and Images Using Deep Features. *18th International Society for Music Information Retrieval Conference (ISMIR 2017)*

# Conversational Systems

# The rise of Conversational Assistants

## Tech Chart of the Day Pandora's new voice search feature knows what you want to hear

The secret sauce for Pandora's Voice Mode is personalized music, regardless of what you ask.



# Task-oriented vs. Chit-Chat

“I want to chat”

“I have a question”

“I need to get this done”

“What should I do?”

Turing Test (talk like a human)

Information consumption

Task completion

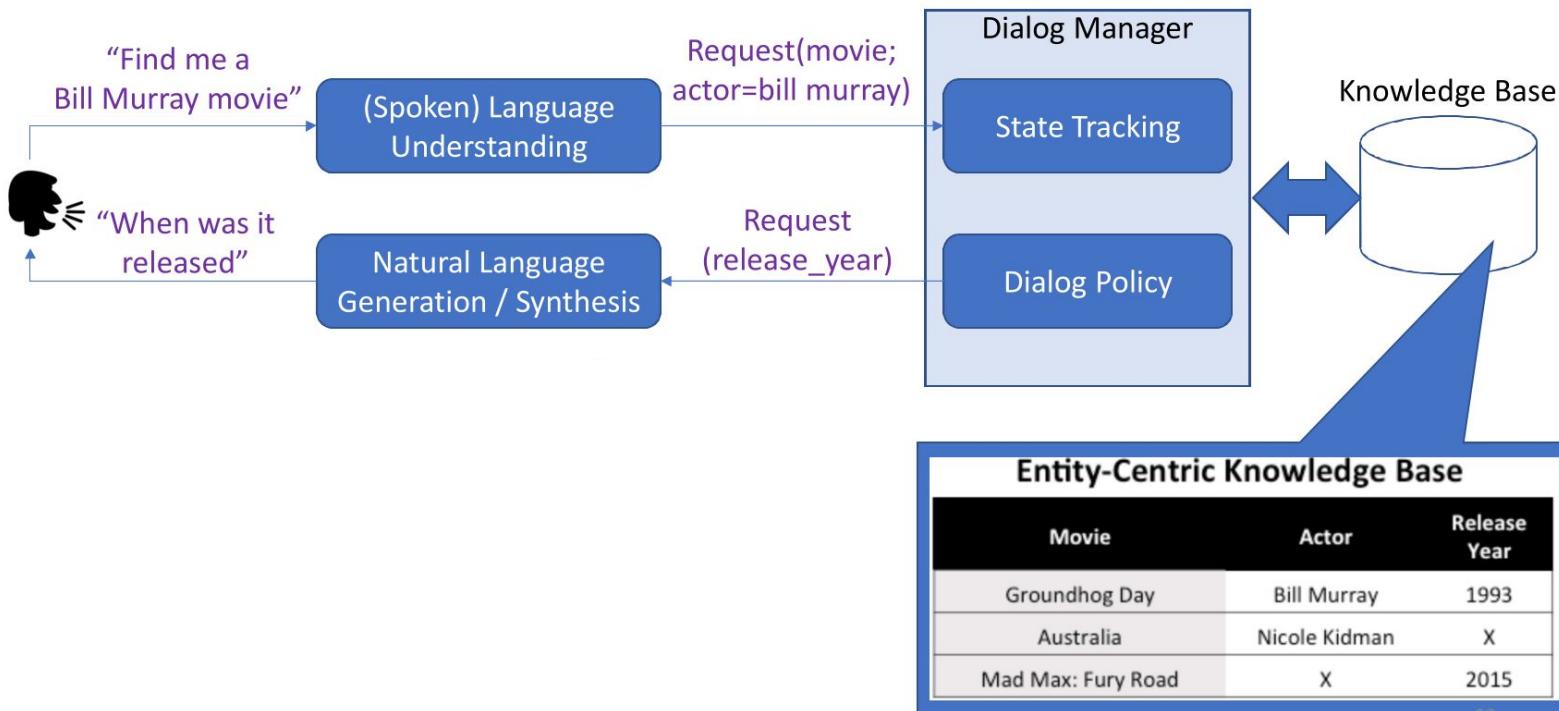
Decision support

Social Chit-Chat

Task-Oriented  
Dialogues



# Task-Oriented Dialog Workflow



# Task-Oriented Dialog Workflow

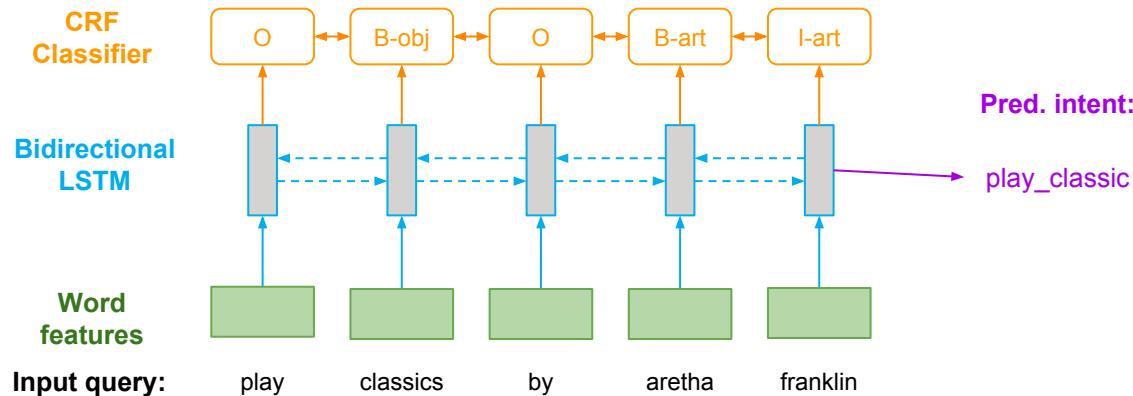
## Spoken Language Understanding

Parse a transcribed oral query into slots and intents

*“Play 90s rock with accordion”*

|               |   |
|---------------|---|
| <i>Slots</i>  | Decade: 90s<br>Genre: rock<br>Instrument: accordion |
| <i>Intent</i> | play_music  |

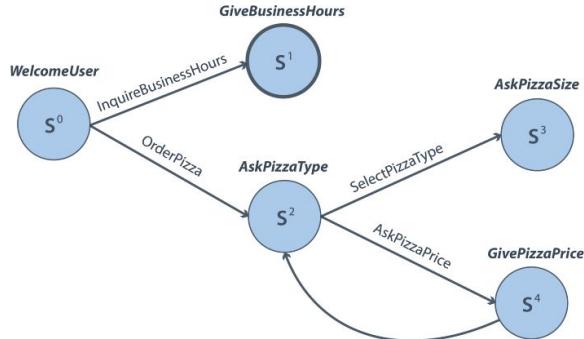
Stacked BiLSTM+CRF ([Reimers and Gurevych, 2017](#))



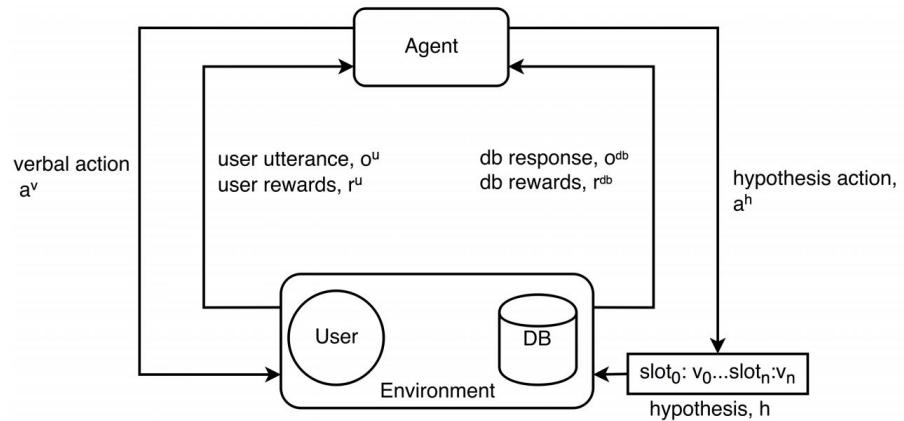
# Task-Oriented Dialog Workflow

## Dialog Policy

Finite State Machines



Reinforcement Learning ([Zhano and Eskenazi, 2016](#))



# Natural Language Generation

Translate a meaning representation into natural language utterances

*Inform(artist=Metallica, genre=heavy\_metal,  
artist\_type=band)*

“Metallica is a heavy metal band”

Template-based NLG

*Inform(artist=\$A, genre=\$G, artist\_type=\$T) ⇒ “\$A is a \$G \$T”*

Generative models

RNN-based NLG ([Wen+, 2015](#))

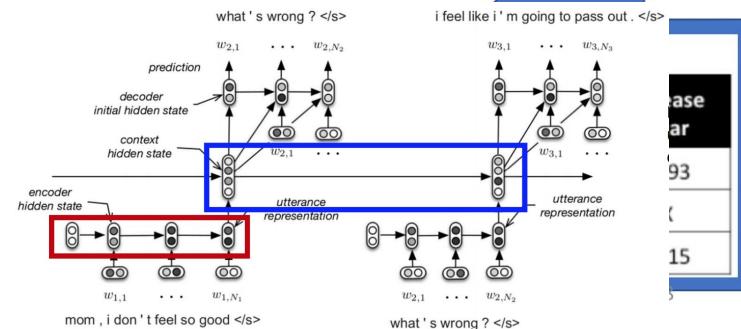


# Data-driven E2E Dialog Systems



Hierarchical Encoder-Decoder ([Serban+, 2016](#))

Memory Networks ([Bordes+, 2017](#))



# Wanna dive more in depth?

## Tutorials

- Neural approaches to Conversational AI [[slides](#), [arxiv](#)]
- Deep learning for Conversational AI [[slides](#)]

## Workshops

- [SCAI](#): Search Oriented Conversational AI (ed. 2019 @ WWW2019 in SF!)
- [CAIR](#): Conversational Approaches to IR

## Blog posts

- [High-density dialog management](#) by b4.ai